

# CSE 127: Introduction to Security

## Lecture 9: Side Channel Attacks

**Nadia Heninger**

Winter 2023

Some material from Dan Boneh, Stefan Savage, Deian Stefan, Keegan Ryan

Last Time: Isolation is key to building secure systems

- Basic idea: protect secrets so that they can't be accessed across a trust boundary
- Mechanisms: privilege separation, least privilege, complete mediation
- **Assumption:** We know what the trust boundaries are and can control access

# How can attackers access protected data?

- Find a bug in an unprotected program
- Find a bug in the kernel, VMM, or runtime system providing protection
- Find a hardware bug that lets you bypass isolation

# The power of abstraction in computer science

"All problems in computer science can be solved by another level of indirection." – David Wheeler

- Computer systems are often built on layers of abstraction
- Physics → hardware → operating system → applications
- An ideal abstraction allows each layer to treat the layer below as a black box with well-defined behavior

# Side channels

Implementations have artifacts and side effects

- How long, how fast, how loud, how hot
- A side channel is a source of information beyond the output specified by an abstraction

# Today

- Overview and history of side channels
- Cache side channels and countermeasures

# Soviet Great Seal Bug

- 1945 Soviet gift to ambassador
- Contained passive listening device
- Would transmit when illuminated at a particular radio frequency
- Designed by Theremin
- Discovered in 1952



[https://en.wikipedia.org/wiki/The\\_Thing\\_\(listening\\_device\)](https://en.wikipedia.org/wiki/The_Thing_(listening_device))

# TEMPEST: US/NATO side channel codename

- WWII: Bell Telephone discovers electromagnetic leakage in one-time pad teleprinters: 100-ft radius
- 1951: CIA rediscovers teleprinter leakage; 200-ft radius
- 1964: TEMPEST shielding rules established



# van Eck Phreaking

"Electromagnetic Radiation from Video Display Units: An Eavesdropping Risk?" Wim van Eck 1985

- 1985: Wim van Eck demonstrates side channel image recovery from CRT monitors with off-the-shelf equipment

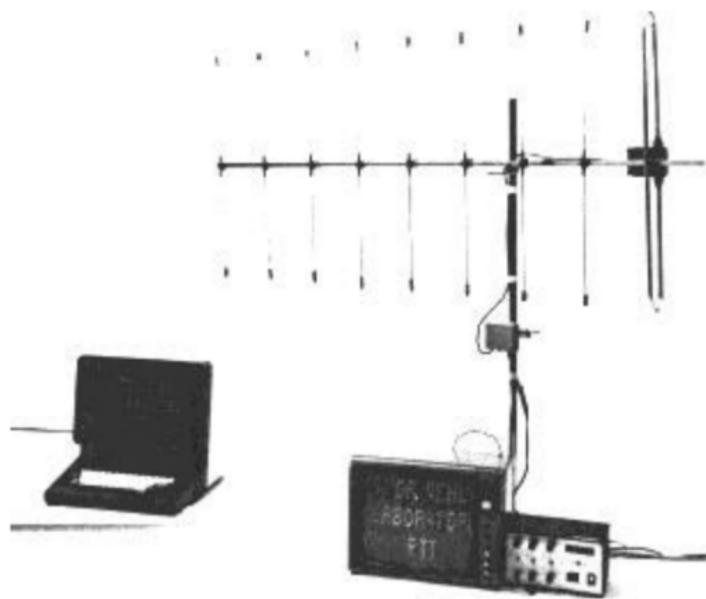


Fig. 1. Eavesdropping set-up using a variable oscillator and a frequency divider to restore synchronization. The picture on the TV is picked up from the radiation of the VDU in the background.

## “Electromagnetic Eavesdropping Risks of Flat-Panel Displays” Kuhn 2004

- Image displays simultaneously along line
  - Pick up radiation from screen connection cable

350 MHz, 50 MHz BW, 12 frames (160 ms) averaged

The quick brown fox jumps over the lazy dog.

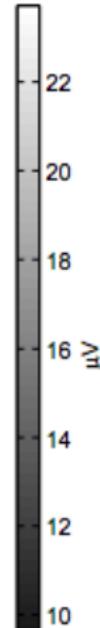
It is well known that electronic equipment produces electromagnetic fields which may cause interference to radio and television reception. The phenomena underlying this have been thoroughly studied over the past few decades. These studies have resulted in internationally agreed methods for measuring the interference produced by equipment. These are needed because the maximum interference levels which equipment may generate have been laid down by law in most countries.

However, interference is not the only problem caused by electromagnetic radiation. It is possible in some cases to obtain information on the signals used inside the equipment when the radiation is picked up and the received signals are decoded. Especially in the case of digital equipment this possibility constitutes a problem, because remote reconstruction of signals inside the equipment may enable reconstruction of the data the equipment is processing.

This problem is not a new one; defence specialists have been aware of it for over twenty years. Information on the way in which this kind of "eavesdropping" can be prevented is not freely available. Equipment designed to protect military information will probably be three or four times more expensive than the equipment likely to be used for processing of non-military information.

(Excerpt from *Klaas van Eck: Electromagnetic Radiation from Video Display Units: An Eavesdropping Risk?* Computers & Security, 4 (1985) 269-286.)

!#\$2A'@\*\*-.#0123456789;.=>@B#CDEFGHJKLNOPQRSTUVWXYZ! - 1  
!abcdeghijklmnopqrstuvwxyz!#\$2A'@\*\*-.#0123456789;.=>? - 1  
@B#CDEFGHJKLNOPQRSTUVWXYZ!# - abcdefghijklmnopqrstuvwxyz!# - 1  
veneck.txt Lines 1-25/26 (END)



# Consumption side channels

How long does this password check take?

```
char pwd[] = "z2n34uzbnqhw4i";  
//...  
  
int check_password(char *buf) {  
    return strcmp(buf, pwd);  
}
```

# Examples of side channels

**Consumption:** How much of a resource is being used to perform an operation?

- Timing
  - Different execution time due to program branches
  - Cache timing attacks
- Power consumption
- Network traffic

**Emission:** What out-of-band signal is generated in the course of performing the operation?

- Electromagnetic radiation
  - Voltage running through a wire produces a magnetic field
- Sound (acoustic attacks)
  - Capacitors discharging can make noises
- Error messages

# Tenex password verification bug

Alan Bell 1974

- Early virtual memory implementation in Tenex computer system.
- Character-at-a-time comparison + interrupt on memory page
- Linear-time password recovery

# "Timing Analysis of Keystrokes and Timing Attacks on SSH"

Song Wagner Tian 2001

- In interactive SSH, keystrokes sent in individual packets
- Build model of inter-keystroke delays by finger, key pair
- Measure packet timing off network, do Viterbi decoding

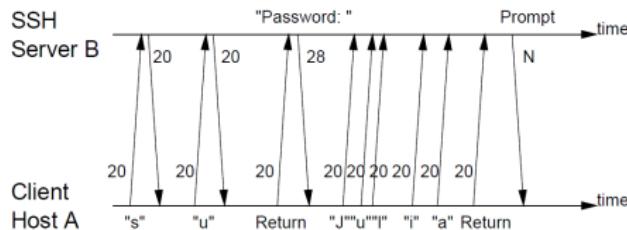


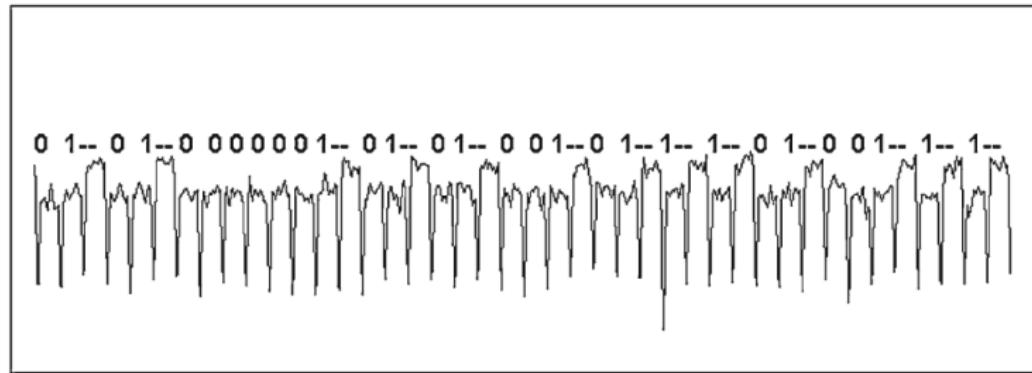
Figure 1: The traffic signature associated with running SU in a SSH session. The numbers in the figure are the size (in bytes) of the corresponding packet payloads.

# Power Analysis Attacks

Kocher Jaffe Jun 98

Side-channel attacks can also leak cryptographic secrets.

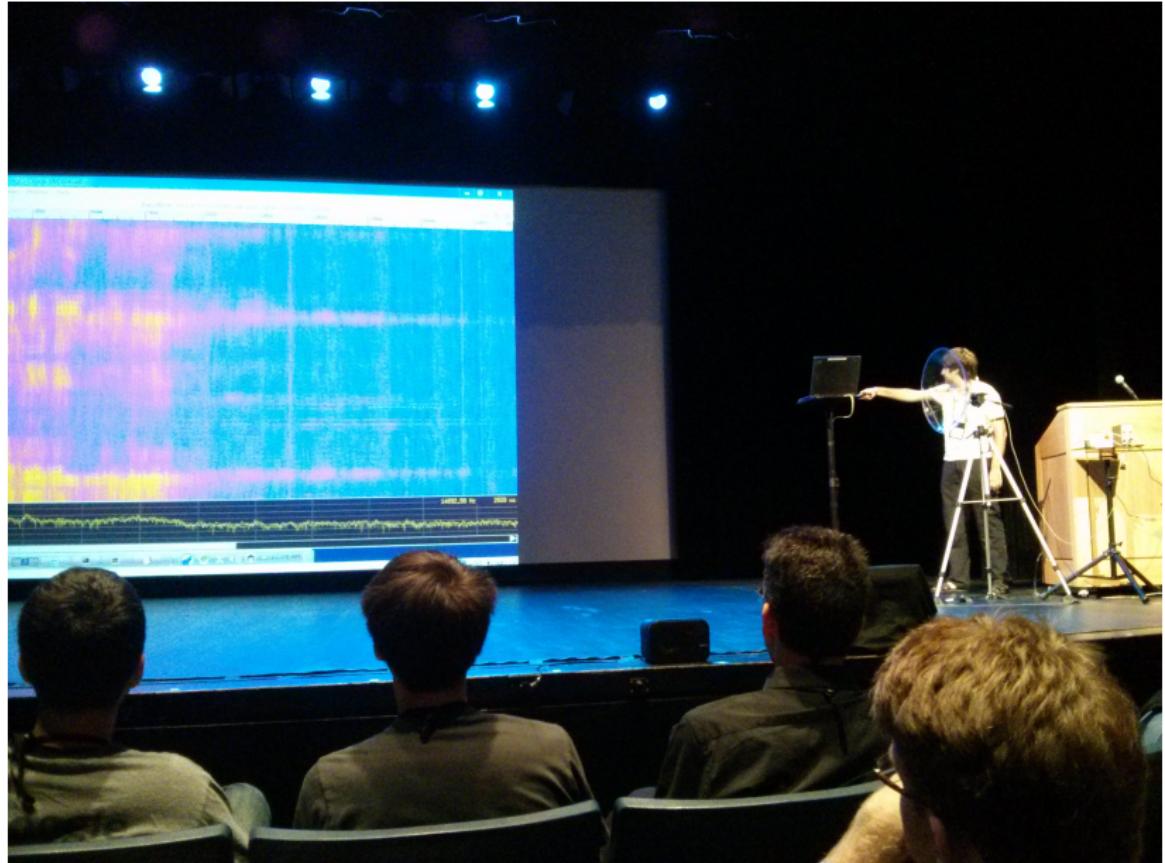
Simple power analysis and differential power analysis  
exploit secret-dependent power consumption.



**Fig. 11** SPA leaks from an RSA implementation

# Acoustic Attacks

Genkin Shamir Tromer 2014



# Browser History Sniffing

Jang, Jhala, Lerner, Shacham 2010

Default web browser behavior: unvisited links are blue and visited links are purple.

Text display attributes available to scripts via DOM.

Victim browser visits malicious website. Malicious website enumerates URLs in invisible portion of site to sniff browser history.

Exploited in the wild.

Fixed in browsers, but surprisingly hard to eliminate all the information leaks.

# Active side channels

- Fault attacks induce computational errors that may leak information
- Attackers can induce faults by:
  - Glitch power, voltage, clock
  - Vary temperature
  - Subject to light, EM radiation

# Using Memory Errors to Attack a Virtual Machine

Govindavajhala Appel 2003

Java heap overflow via glitched address of function pointer.



**Figure 3. Experimental setup to induce memory errors, showing a PC built from surplus components, clip-on gooseneck lamp, 50-watt spotlight bulb, and digital thermometer. Not shown is the variable AC power supply for the lamp.**

# Types of RAM

- Volatile memory: Data retained only as long as power is on
- Persistent memory like flash or magnetic disks retains data without power

## SRAM

- SRAM retains bit value as long as power is on without any refresh
- Faster, lower density, higher cost
- Has a “burn-in” phenomenon where on startup it tends to flip bit to “remembered bit”

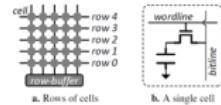
## DRAM

- DRAM requires periodic refresh to retain stored value
- Capacitors charged to store data
- Higher density, lowered cost
- “Cold boot attacks” exploited capacitor discharge time to read sensitive data from physical memory

# Rowhammer attacks

Seaborn and Dullien 2015

- DRAM cells are grouped into rows
- All cells in a row are refreshed together
  - Repeatedly opening and closing a row within a refresh interval causes disturbance errors in adjacent rows.
  - Attacker running attack process on same machine as victim can cause bits to flip in victim's memory



## Aside: Covert Channels

Side channels are inadvertent artifacts of the implementation

- Covert channels: same idea, but on purpose
- One party is trying to leak information across a trust boundary by encoding that information into some side channel
- Examples: variation in time, memory usage
- Difficult to protect against

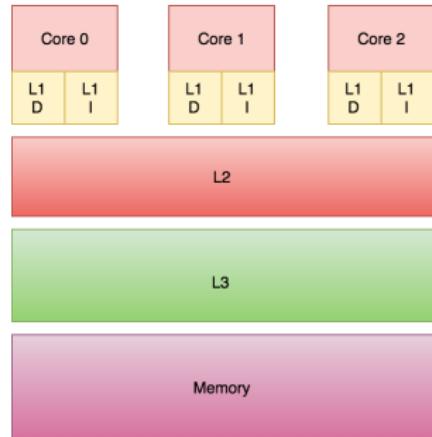
# Radio covert channels

Guri et al. 2015

- CPU memory bus emits radio waves on GSM frequencies
- Covert channel to hop air gap: Software performs memory reads; signal is detected by cell phone

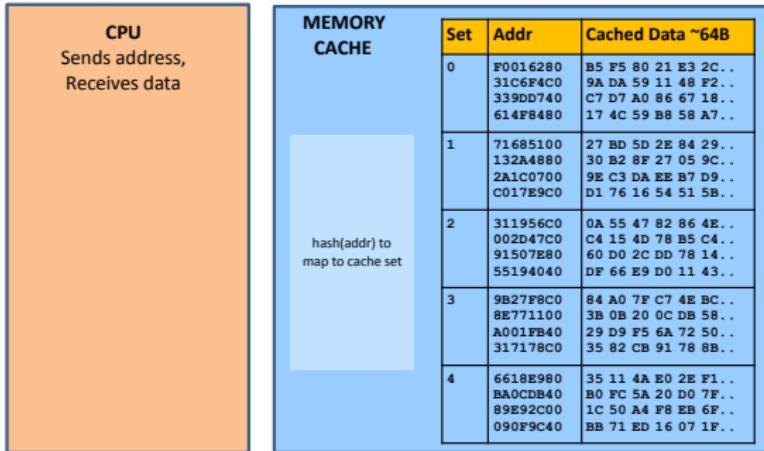
# Memory and cache

- Main memory is large and slow
- Processors have faster, smaller caches to store more recently used memory closer to cores
- Caches organized in hierarchy: closer to the core are faster and smaller



# Memory and cache

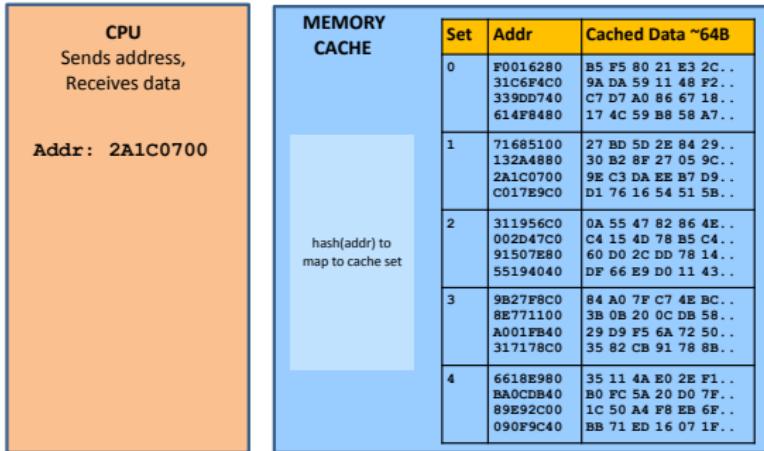
Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



**MAIN  
MEMORY**  
Big, slow  
e.g. 16GB SDRAM

# Memory and cache

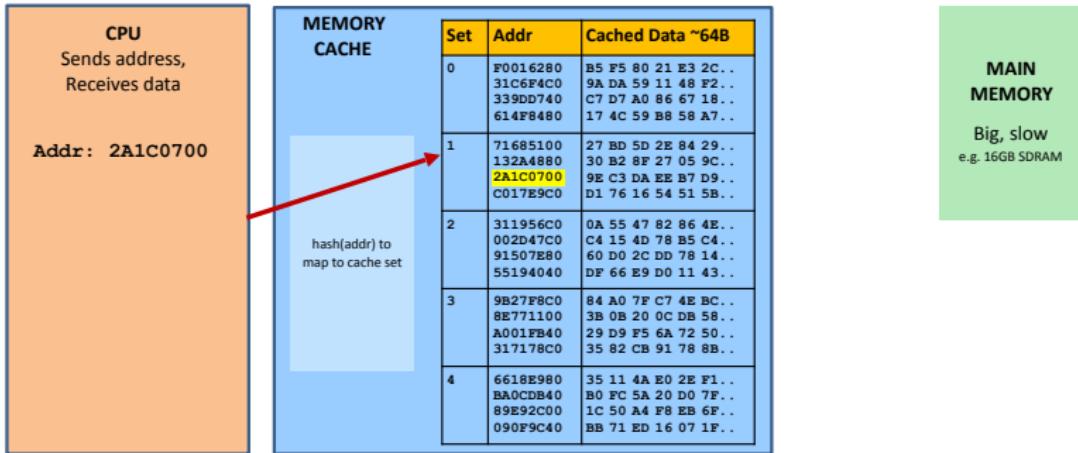
Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



**MAIN MEMORY**  
Big, slow  
e.g. 16GB SDRAM

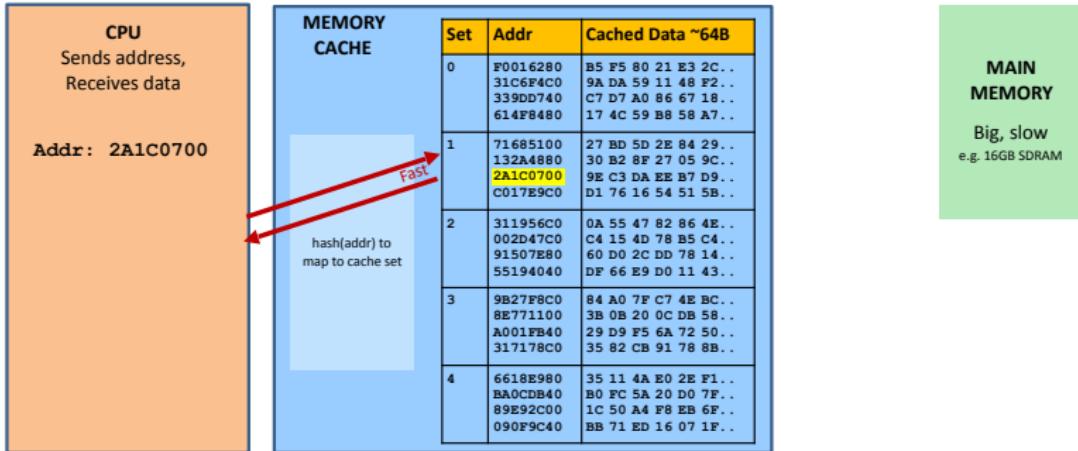
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



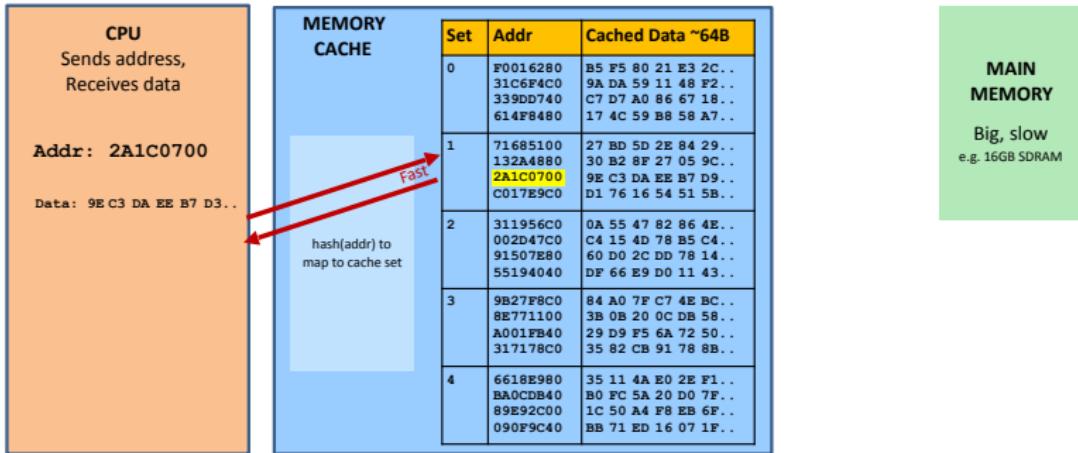
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



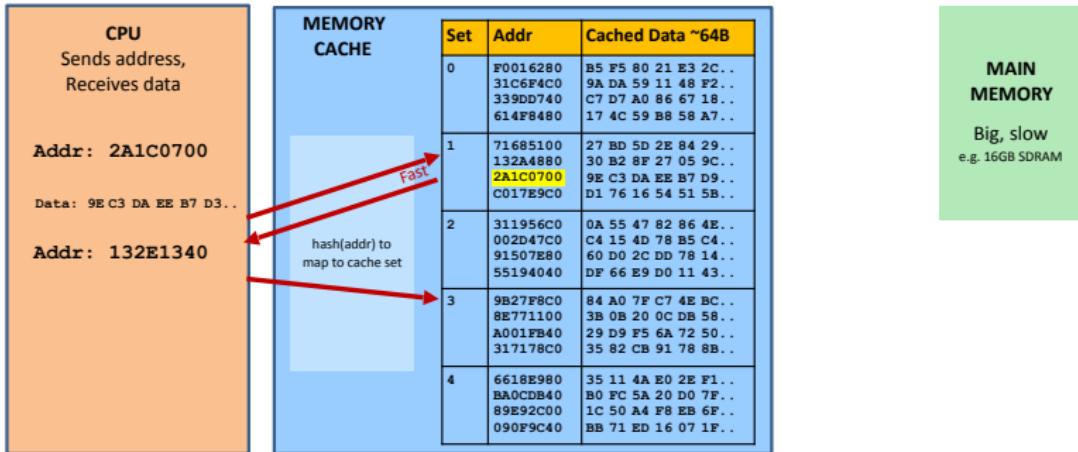
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



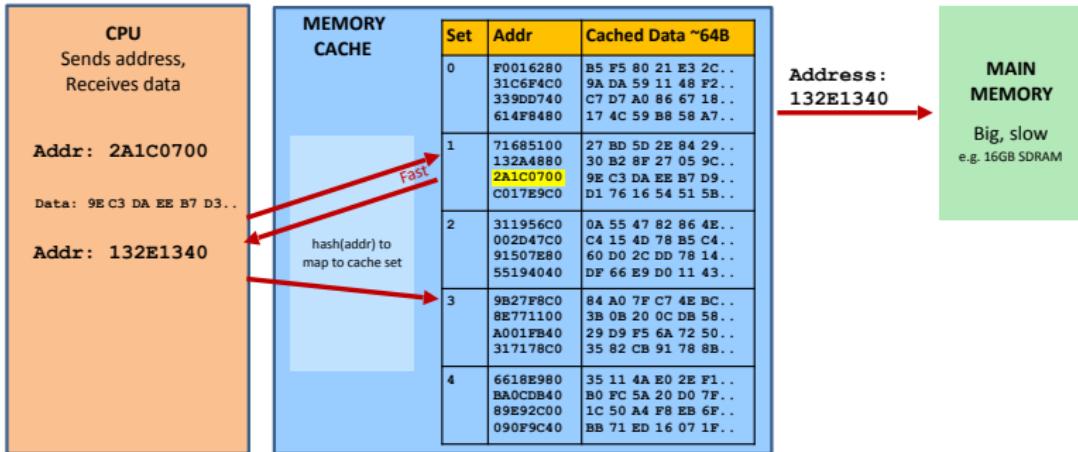
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



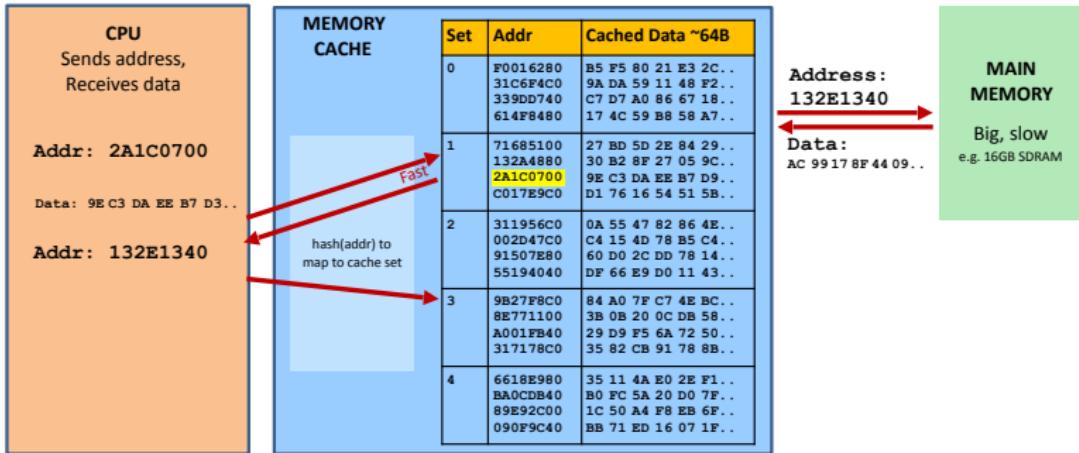
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



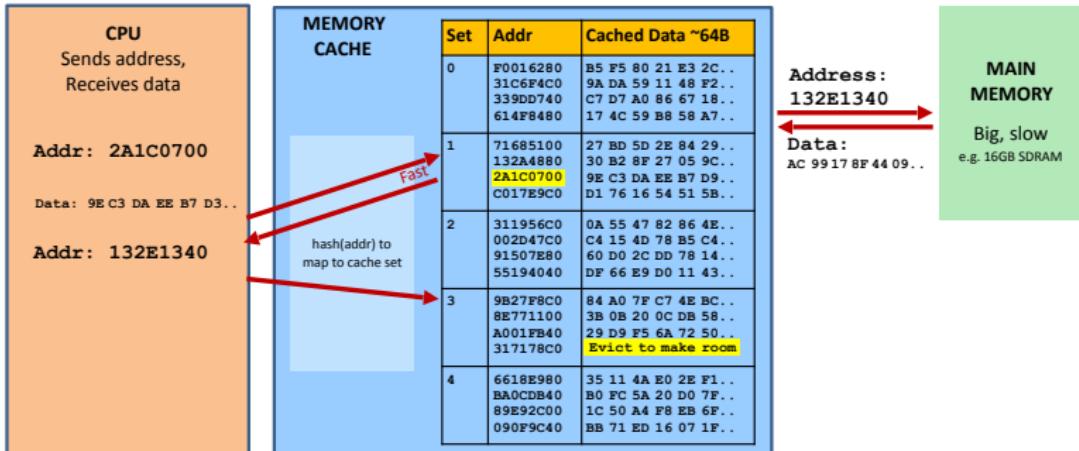
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



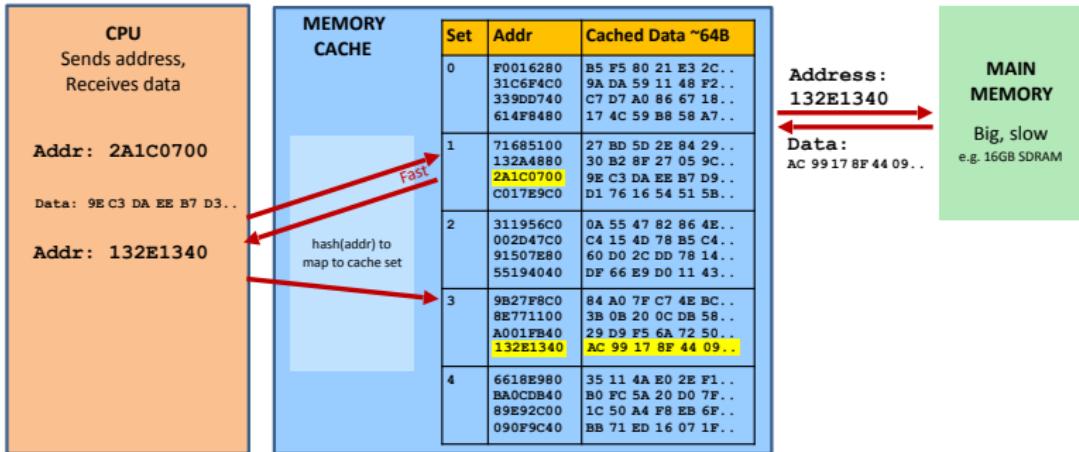
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



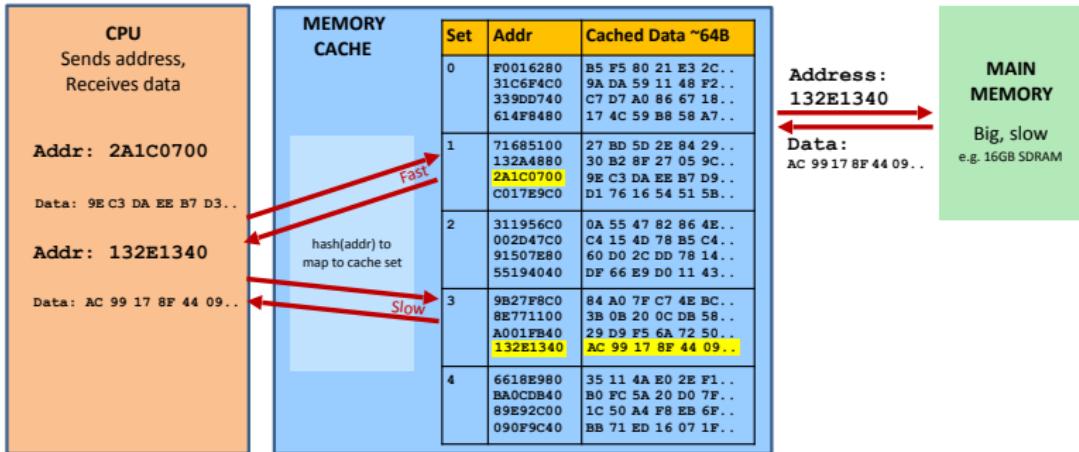
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



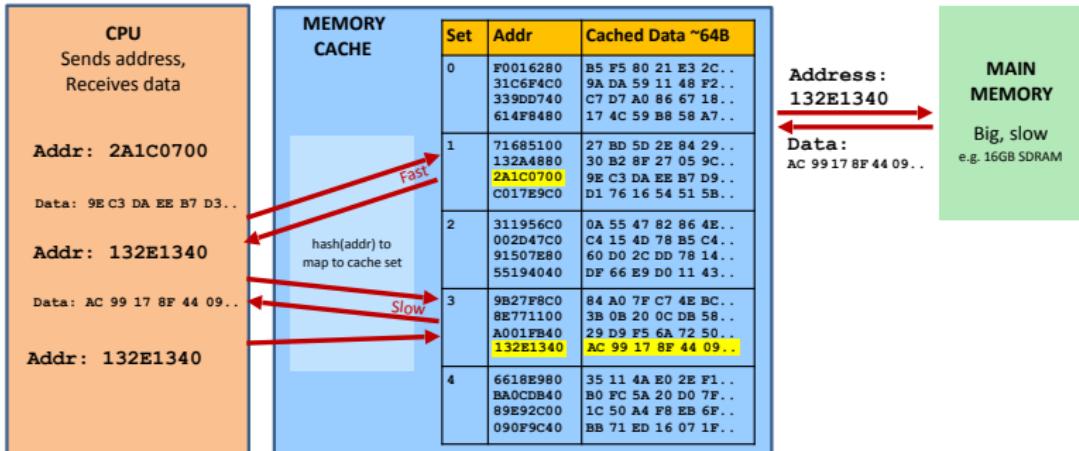
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



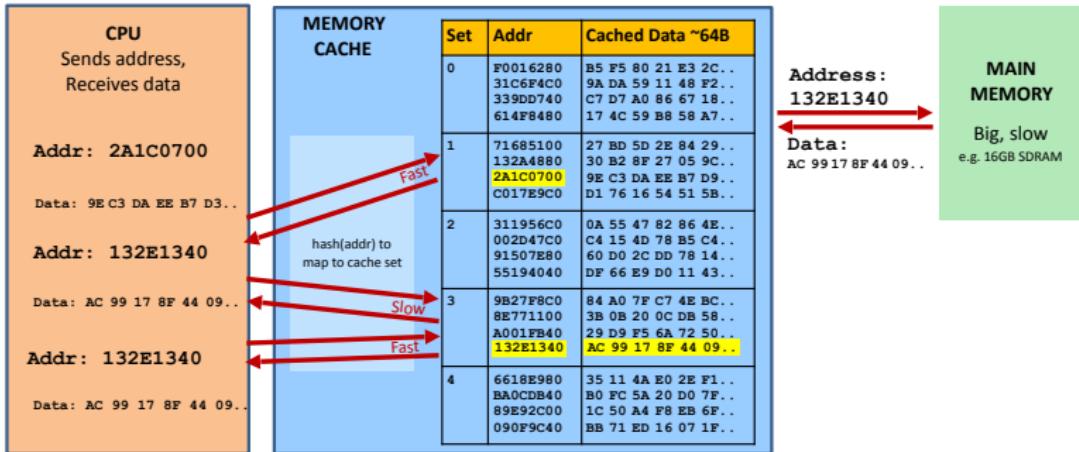
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



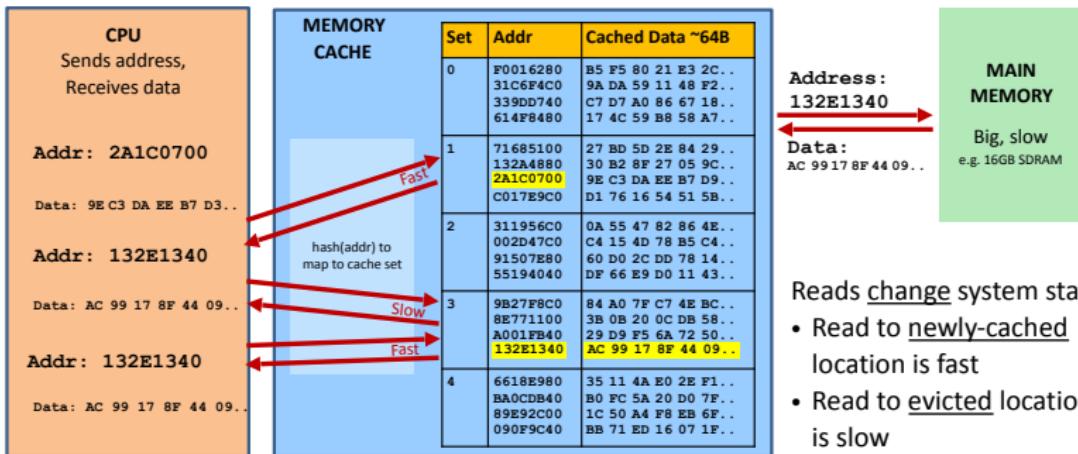
# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



# Memory and cache

Caches hold local (fast) copy of recently-accessed 64-byte chunks of memory



# Cache timing side channel attacks

- Caches are a shared system resource
- Not isolated by process, VM, or privilege level
- An attacker who can run code on same physical hardware can abuse this shared resource to learn information from another process or VM

# Cache timing attack threat model

- Attacker and victim are isolated (separate processes) on same physical system
- Attacker is able to invoke (directly or indirectly) functionality exposed by victim
  - Examples?
- Attacker does not have direct access to contents of victim memory
- Attacker will use shared cache access to infer information about victim memory contents

# Cache timing attack vector

- Many algorithms have memory access patterns that are dependent on sensitive memory contents
  - Examples?
- If attacker can observe access patterns they can learn secrets

# Cache timing attack options

- **Prime:** Place a known address in the cache by reading it
- **Evict:** Access memory until address is no longer cached (force capacity misses)
- **Flush:** Remove an address from the cache (`clflush` on x86)
- **Measure:** Precisely (down to the cycle) how long it takes to do something (`rdtsc` on x86)
- Attack form: Manipulate cache into known state, make victim run, infer what changed after run

# Three basic techniques

- Evict and time
  - Evict things from the cache and measure if victim slows down as a result
- Prime and probe
  - Place things in the cache, run the victim, and see if you slow down as result
- Flush and reload
  - Flush a particular line from the cache, run the victim, and see if your accesses are still fast

# Mitigating side channels

- Use constant-time programming techniques
- Eliminate secret-dependent execution or branches
- Hide/blind inputs

# Constant-time programming: Which is faster?

```
void foo(double x) {  
    double z, y = 1.0;  
    for (uint32_t i = 0; i < 100000000; i++) {  
        z = y*x;  
    }  
}
```

- a. `foo(1.0);`
- b. `foo(1.0e-323);`
- c. They are the same

# Constant-time programming: Which is faster?

```
void foo(double x) {  
    double z, y = 1.0;  
    for (uint32_t i = 0; i < 100000000; i++) {  
        z = y*x;  
    }  
}
```

- a. `foo(1.0);` ←
- b. `foo(1.0e-323);`
- c. They are the same

Some instructions take different amounts of time depending on operands.

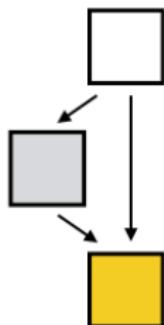
Avoid variable-time instructions.

# Control flow introduces time variability

```
m=1
for i = 0 ... len(d):
    if d[i] = 1:
        m = c * m mod N
        m = square(m) mod N
return m
```

# If-statements on secrets are unsafe

```
s0;  
if (secret) {  
    s1;  
    s2;  
}  
s3;
```



secret	run	
true	s0;s1;s2;s3;	4
false	s0;s3;	2

# Does padding the else branch work?

```
if (secret) {  
    s1;  
    s2;  
} else {  
    s1';  
    s2';  
}
```

where s1 and s1' take same amount of time

# Does padding the else branch work?

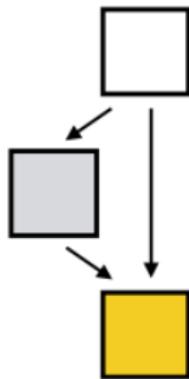
```
if (secret) {  
    s1;  
    s2;  
} else {  
    s1';  
    s2';  
}
```

where s1 and s1' take same amount of time

- **Problem:** Instructions are loaded from cache  
    ⇒ observable
- **Problem:** Hardware tries to predict where branch goes  
    ⇒ success or failure of prediction is observable

# Constant-time conditional arithmetic assignment

```
if (secret) {  
    x = a;  
}  
  
x = secret * a  
+ (1-secret) * x
```



# Constant-time conditional arithmetic assignment

```
if (secret) {  
    x = a;           x = secret * a + (1-secret) * x  
} else {  
    x = b;           x = (1-secret) * b + secret * x  
}
```

Writing constant-time code is hard.

There are tools to help but most code is still written by hand.

It can be slower, larger, and more complex.

(This is one of the reasons people recommend not doing your own cryptographic implementations.)

## Speculative Execution

- CPUs can guess likely program path and do speculative execution
- Example

```
if (uncached_value == 1) // load from memory  
    a = compute(b)
```

- Branch predictor guesses if() is true based on prior history
- Starts executing compute(b) speculatively
- When value arrives from memory, check if guess was correct:
  - Correct: Save speculative work → performance gain
  - Incorrect: Discard speculative work → no harm?

# Spectre and Meltdown

Lipp et al., Kocher et al. 2017

## Misspeculation

- Exceptions and incorrect branch prediction can cause “rollback” of transient instructions
- Old register states are preserved, can be restored
- Memory writes are buffered, can be discarded
- Cache modifications are not restored



- Spectre and Meltdown carry out cache attacks against speculatively loaded data so that an unprivileged attacker process can read kernel memory, break ASLR, etc.

# Today

- Overview and history of side channels
- Cache side channels and countermeasures