



CSE 127: Computer Security

Stack Buffer Overflows

Deian Stefan

Adopted from Kirill Levchenko, Stefan Savage, and Hovav Shacham

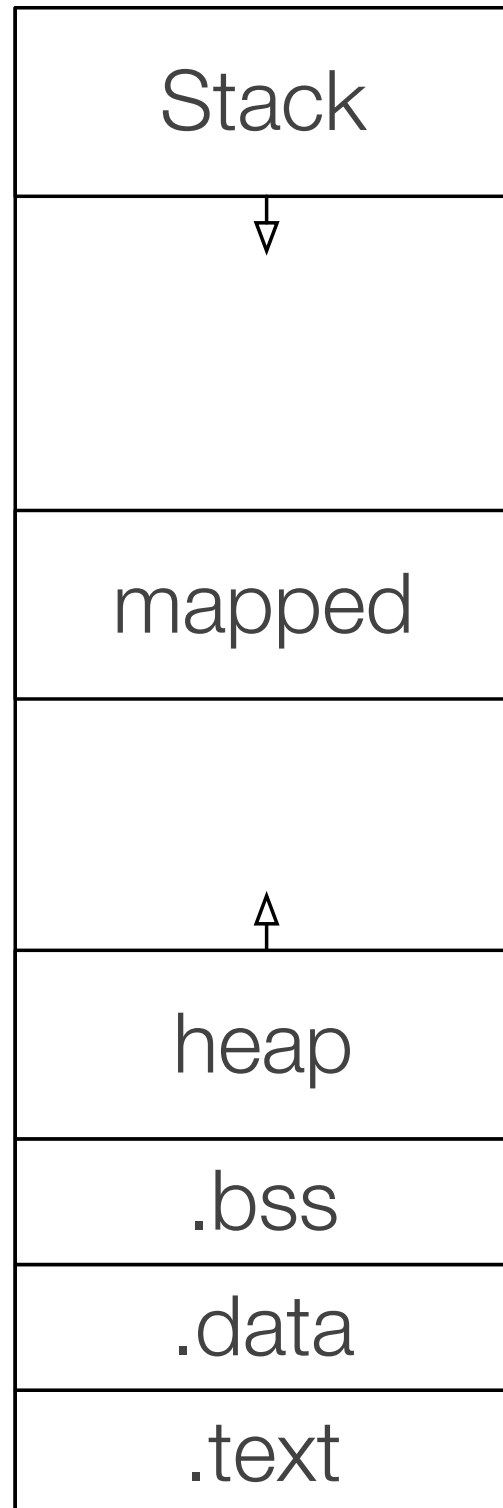
Control Flow Hijacking Defenses

- Avoid unsafe functions
- Stack canary
- Separate control stack
- Address Space Layout Randomization (ASLR)
- Memory writable or executable, not both (W^X)
- Control flow integrity (CFI)

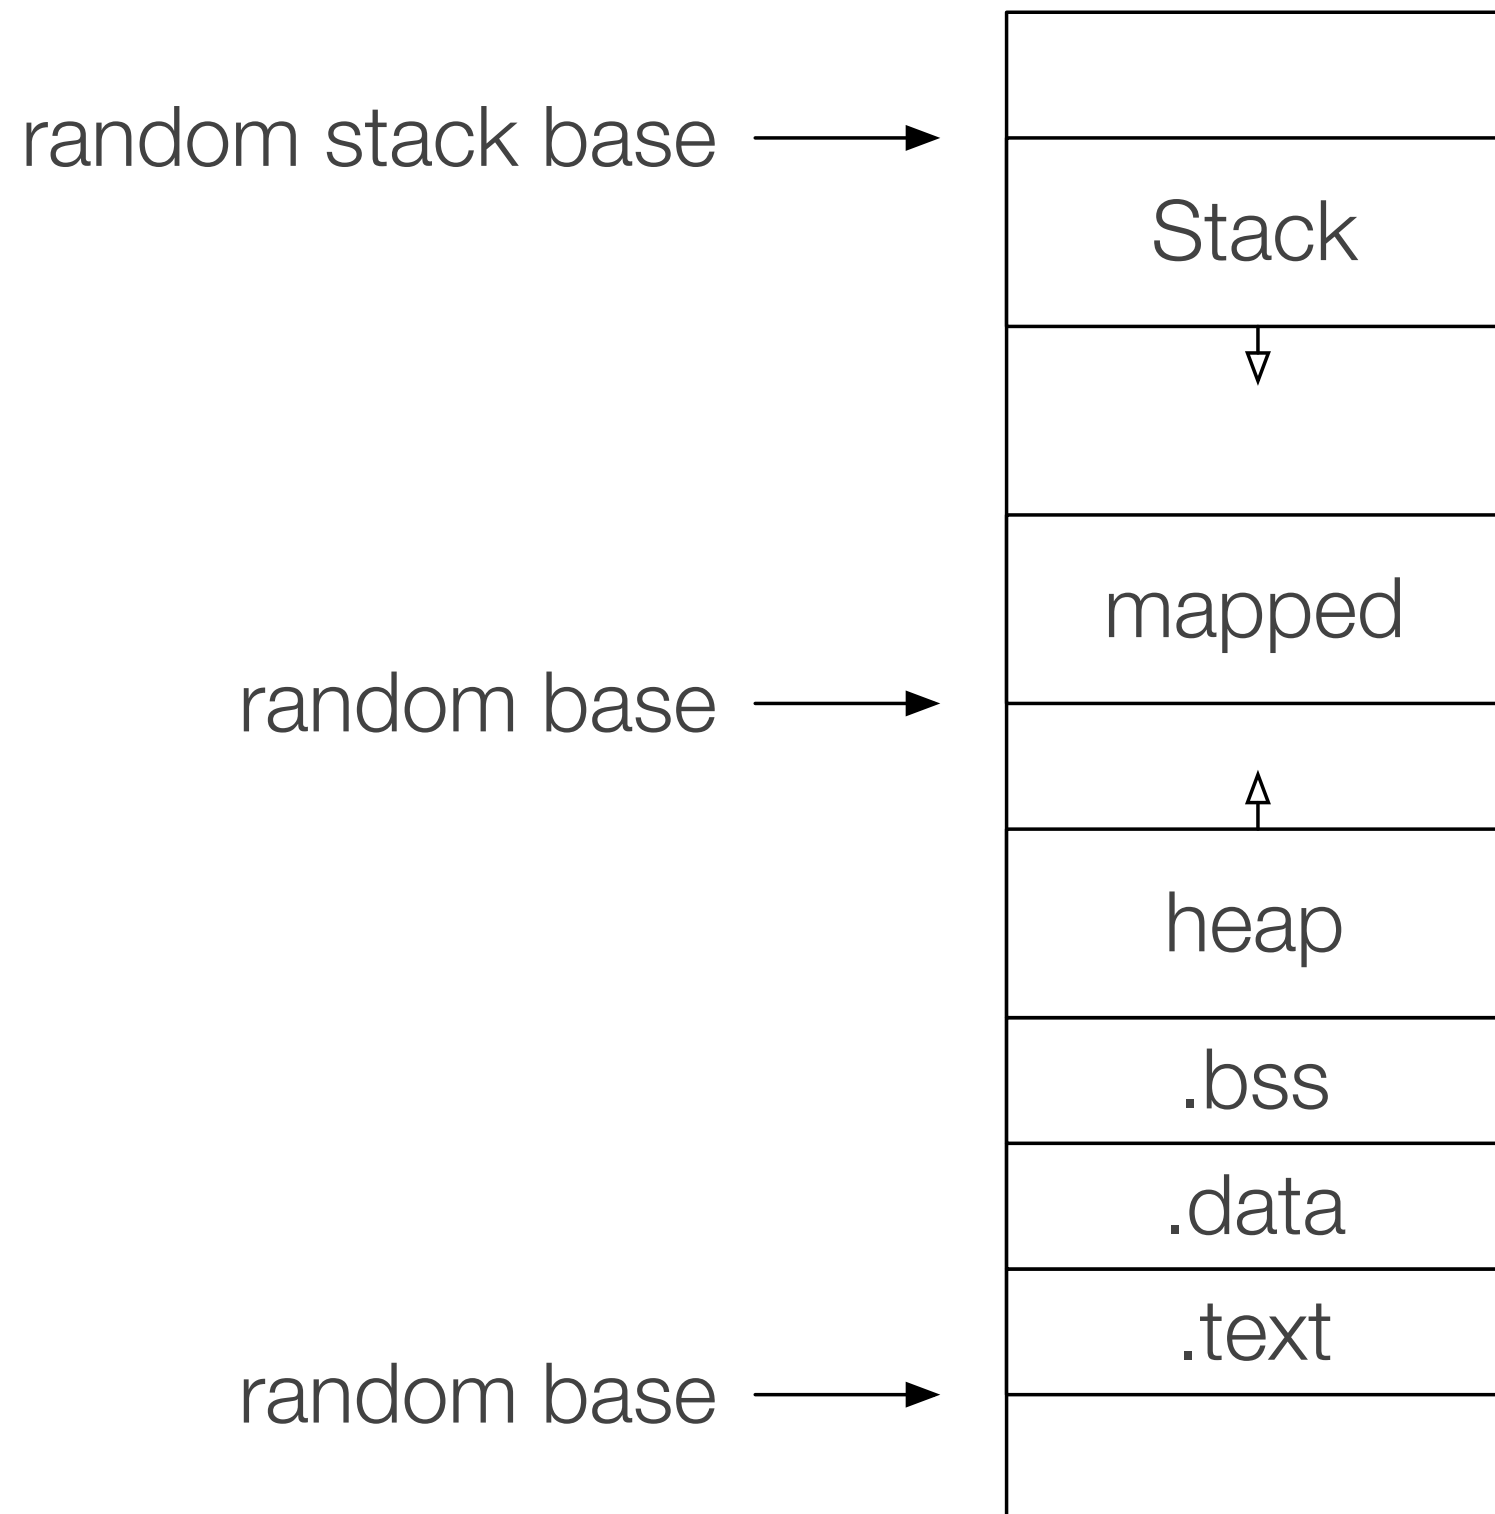
Address Space Layout Randomization

- Change location of stack, heap, code, static vars
- Works because attacker needs address of shellcode
- Layout must be unknown to attacker
 - Randomize on every launch (best)
 - Randomize at compile time
- Implemented on most modern OSes in some form

Traditional Memory Layout




PaX Memory Layout




32-bit PaX ASLR (x86)


Stack:



Mapped area:



Executable code, static variables, and heap:



On the Effectiveness of Address-Space Randomization

Hovav Shacham
Stanford University

hovav@cs.stanford.edu

Matthew Page
Stanford University

mpage@stanford.edu

Ben Pfaff
Stanford University

blp@cs.stanford.edu

Eu-Jin Goh
Stanford University

eujin@cs.stanford.edu

Nagendra Modadugu
Stanford University

nagendra@cs.stanford.edu

Dan Boneh
Stanford University

dabo@cs.stanford.edu

Derandomizing ALSR

- **Attack goal:** call `system()` with attacker arg
- **Target:** Apache daemon
 - **Vulnerability:** buffer overflow in `ap_getline()`

```
char buf[64];
```

```
...
```

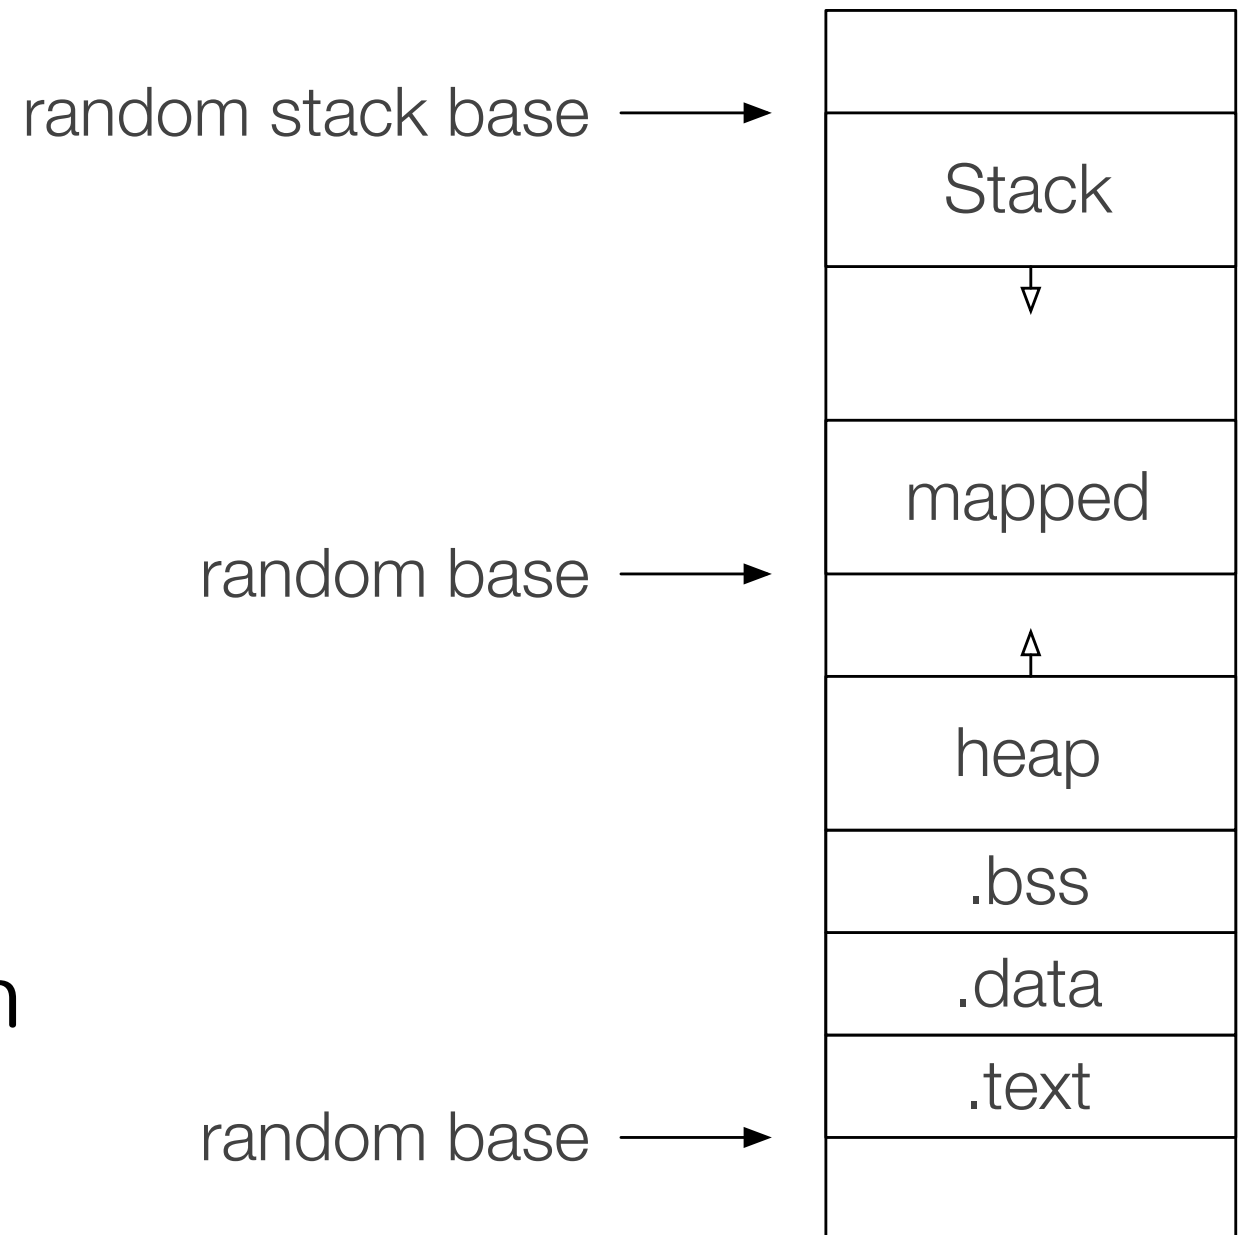
```
strcpy(buf, s); // overflow
```


Defense assumptions

- W^X enabled
- PaX ASLR enabled
 - Apache forks child processes to handle client interaction
 - How does re-randomization work?

Planning the Attack


- Can we inject shell code on the stack?
 - A: yes, **B: no**
- Call system in libc
 - Located in mapped region



Derandomizing ASLR

- **Stage 1: Find base of mapped region**

Mapped area:



- **Stage 2: Call system() with command string**

Finding base of mapped region

- Overflow buffer in `ap_getline()`
- Overwrite saved EIP with guessed location of `usleep`

...
<code>ap_getline()</code> args
saved EIP
saved EBP
buffer to overflow

Finding base of mapped region

- Overflow buffer in `ap_getline()`
- Overwrite saved EIP with guessed location of `usleep`
 - Base + offset of `usleep` in mapped region
- Provide non-zero byte argument to `usleep()`

...
<code>ap_getline()</code> args
saved EIP
saved EBP
buffer to overflow

Finding base of mapped region

- If we guessed `usleep()` address right
 -
- If we guessed `usleep()` address wrong
 -
- Use this to tell if we guessed base of mapped region correctly

Finding base of mapped region

- If we guessed `usleep()` address right
 - Server will freeze for 16 seconds, then crash
- If we guessed `usleep()` address wrong
 -
- Use this to tell if we guessed base of mapped region correctly

Finding base of mapped region

- If we guessed `usleep()` address right
 - Server will freeze for 16 seconds, then crash
- If we guessed `usleep()` address wrong
 - Server will (likely) crash immediately
- Use this to tell if we guessed base of mapped region correctly

Finding base of mapped region

```
ap_getline:
```

```
...
```

```
pop ebp
```

```
ret
```

```
usleep:
```

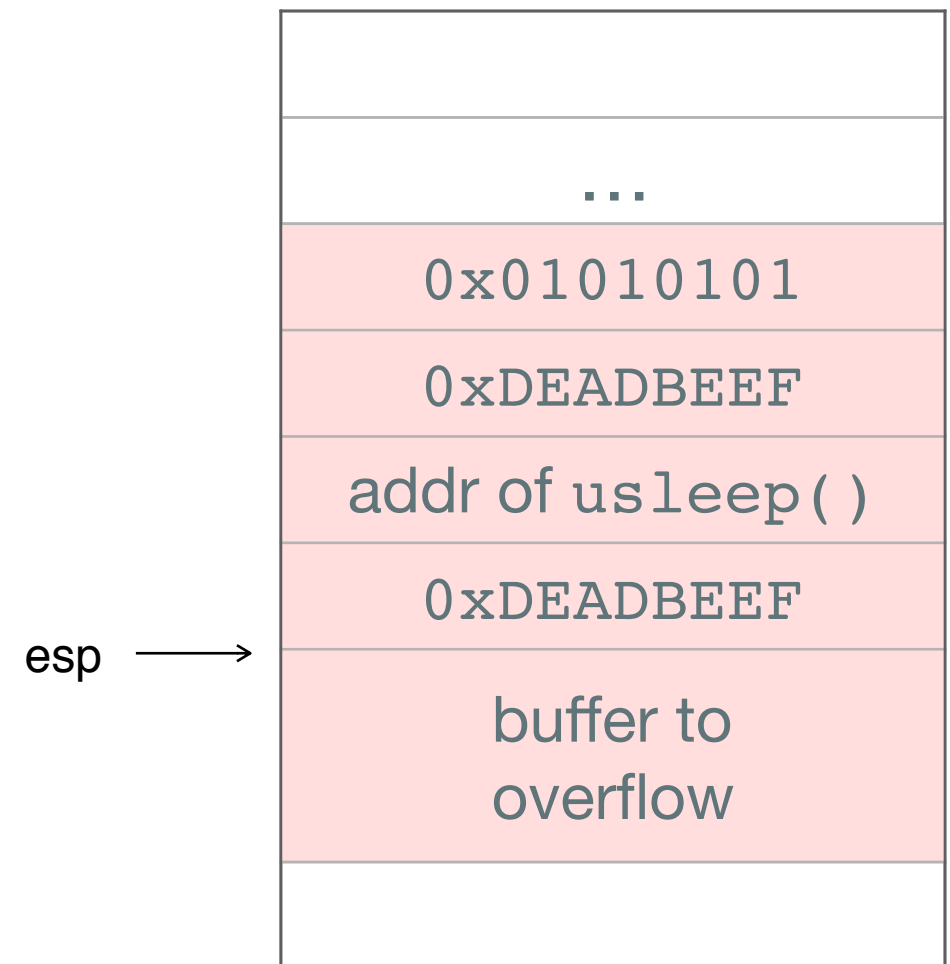
```
...
```

```
ret
```

...
ap_getline() args
saved EIP
saved EBP
buffer to overflow

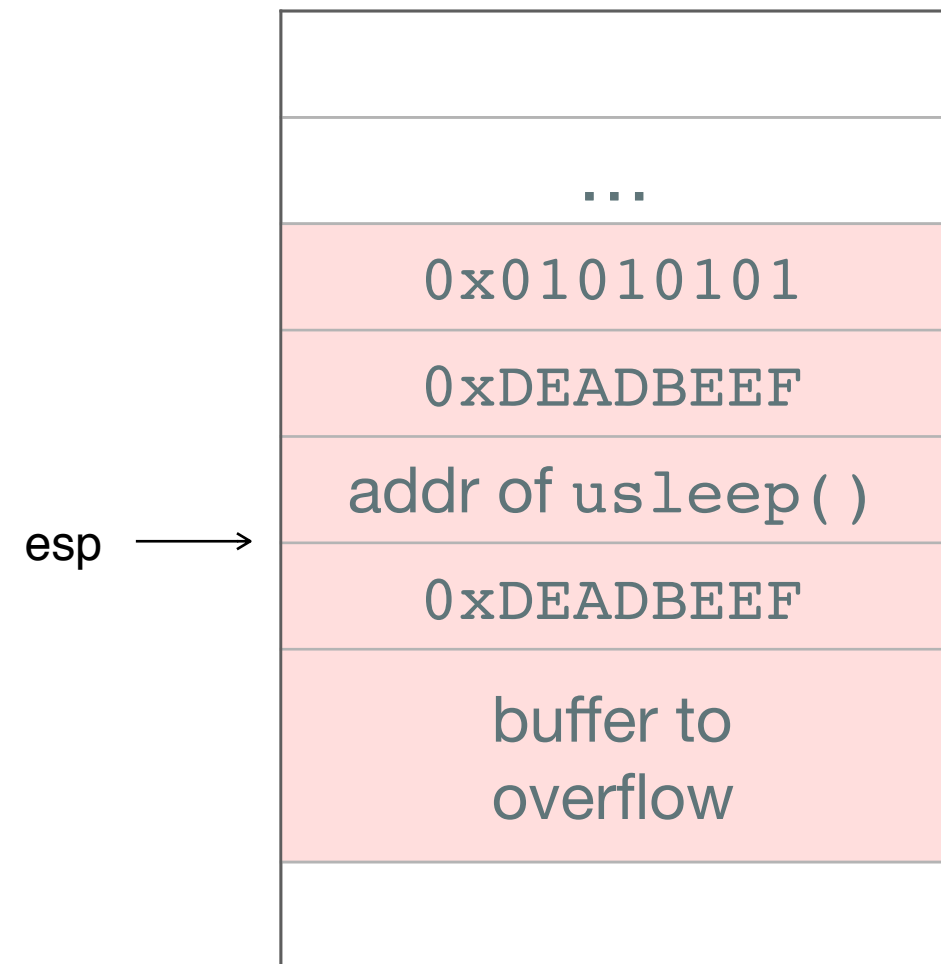
Finding base of mapped region

```
ap_getline:  
...  
eip → pop ebp  
      ret  
  
usleep:  
...  
ret
```



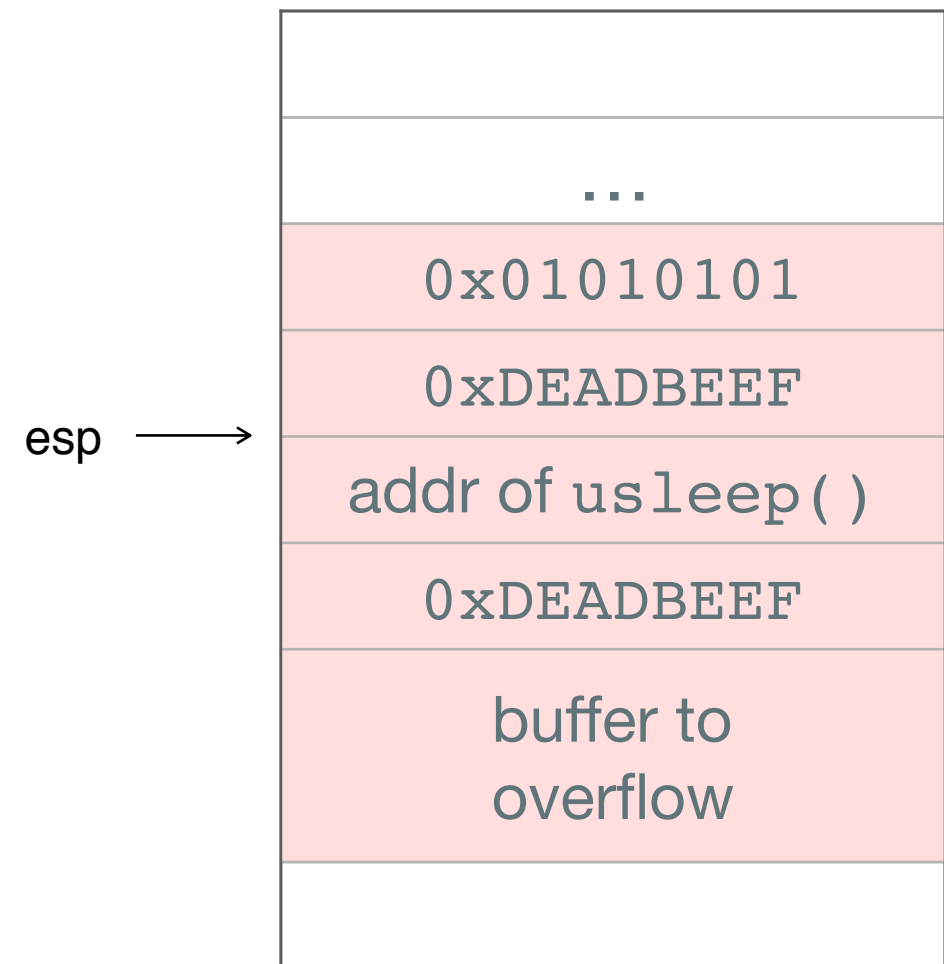
Finding base of mapped region

```
ap_getline:  
...  
eip → pop ebp  
      ret  
  
usleep:  
...  
ret
```



Finding base of mapped region

```
ap_getline:  
...  
pop ebp  
eip → ret  
  
usleep:  
...  
ret
```



Finding base of mapped region

```
ap_getline:
```

```
...
```

```
pop ebp
```

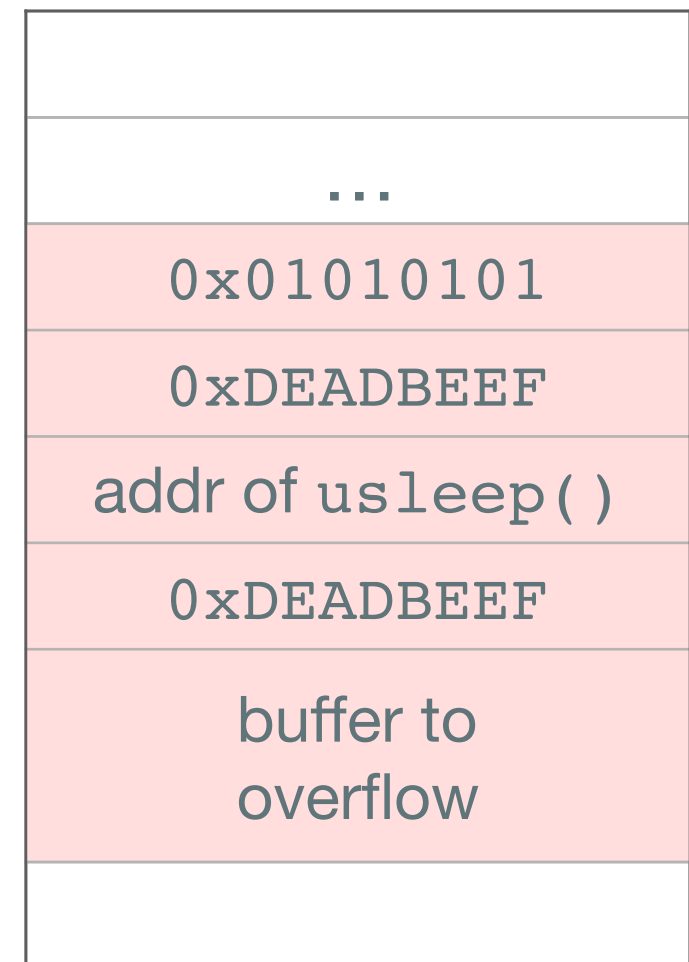
```
ret
```

```
eip → usleep:
```

```
...
```

```
ret
```

esp →



Finding base of mapped region

```
ap_getline:
```

```
...
```

```
pop ebp
```

```
ret
```

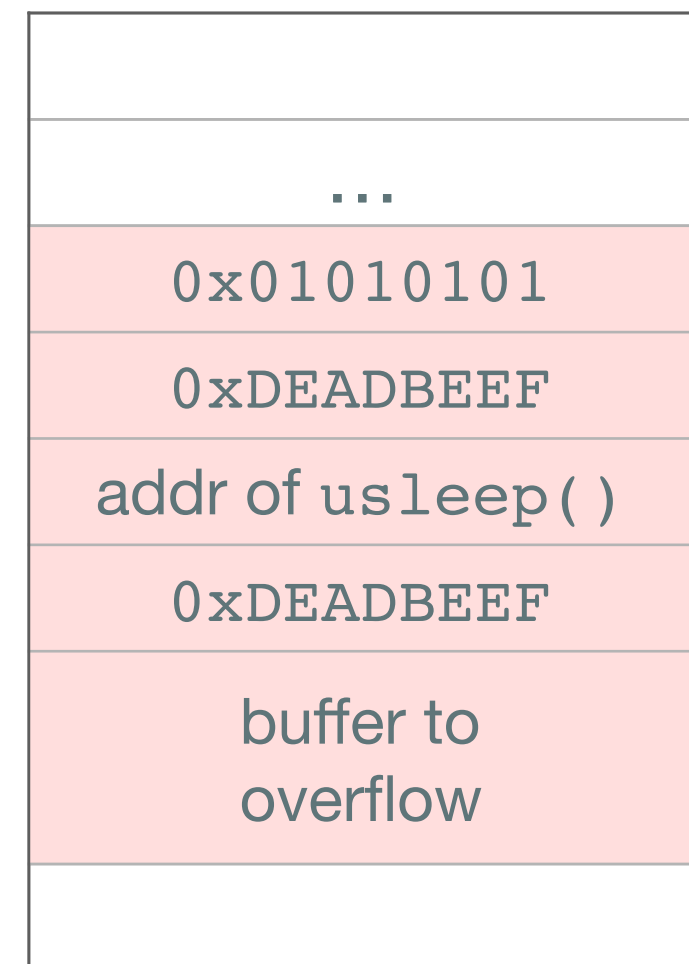
```
eip → usleep:
```

```
...
```

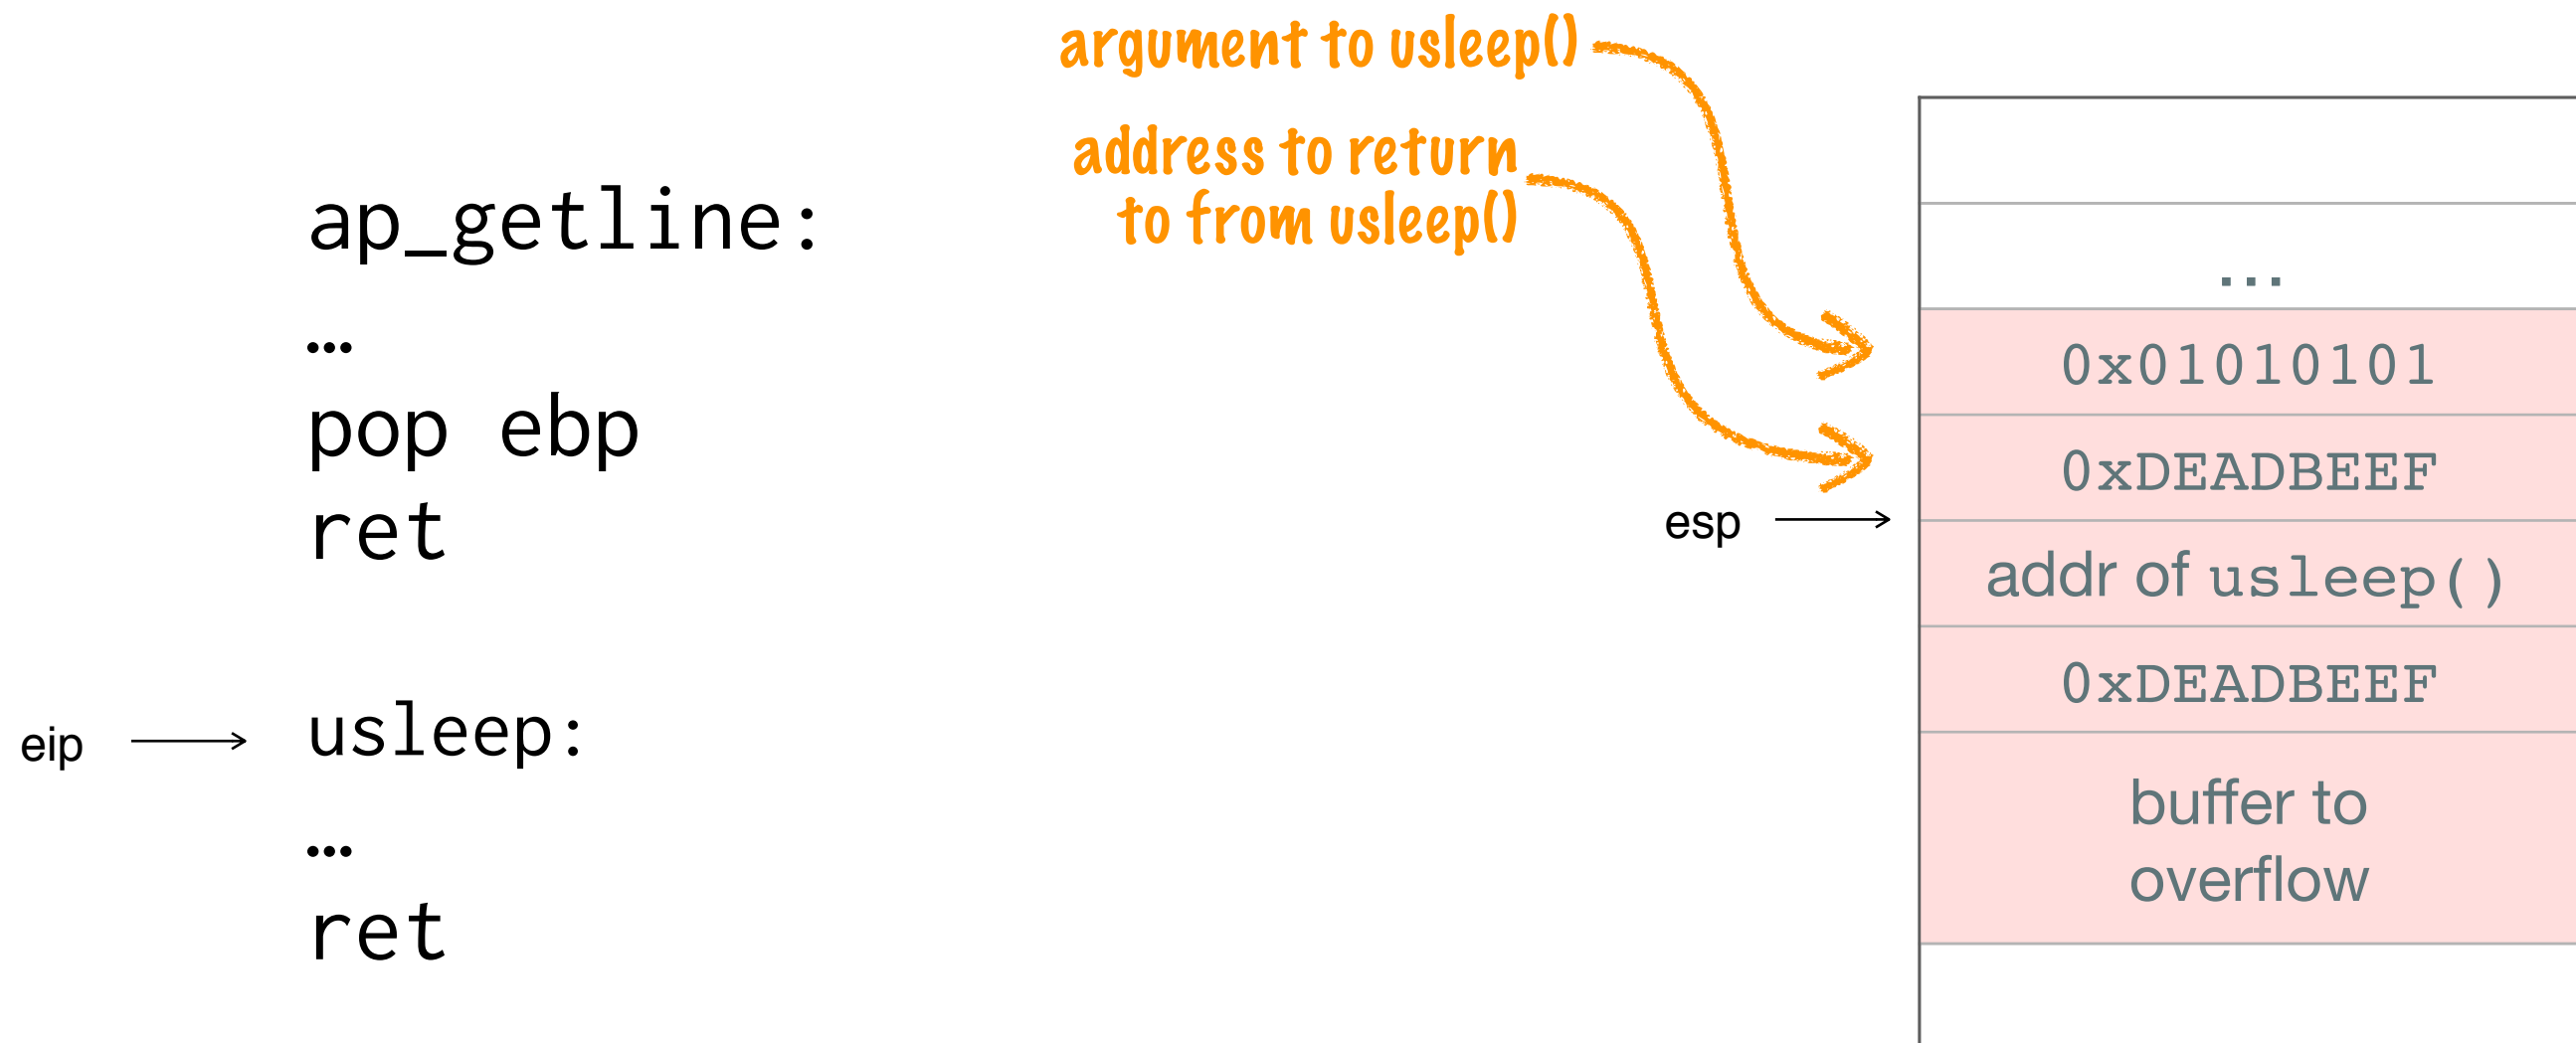
```
ret
```

argument to usleep()

esp →



Finding base of mapped region



Finding base of mapped region

```
ap_getline:
```

```
...
```

```
pop ebp
```

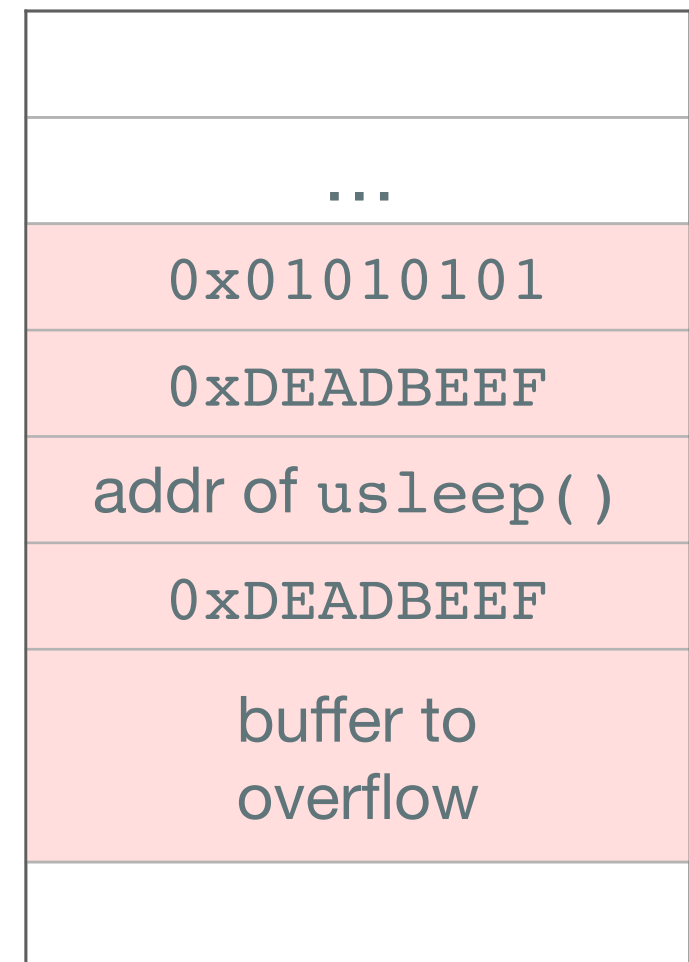
```
ret
```

```
usleep:
```

```
eip → ...
```

```
ret
```

esp →



Finding base of mapped region

```
ap_getline:
```

```
...
```

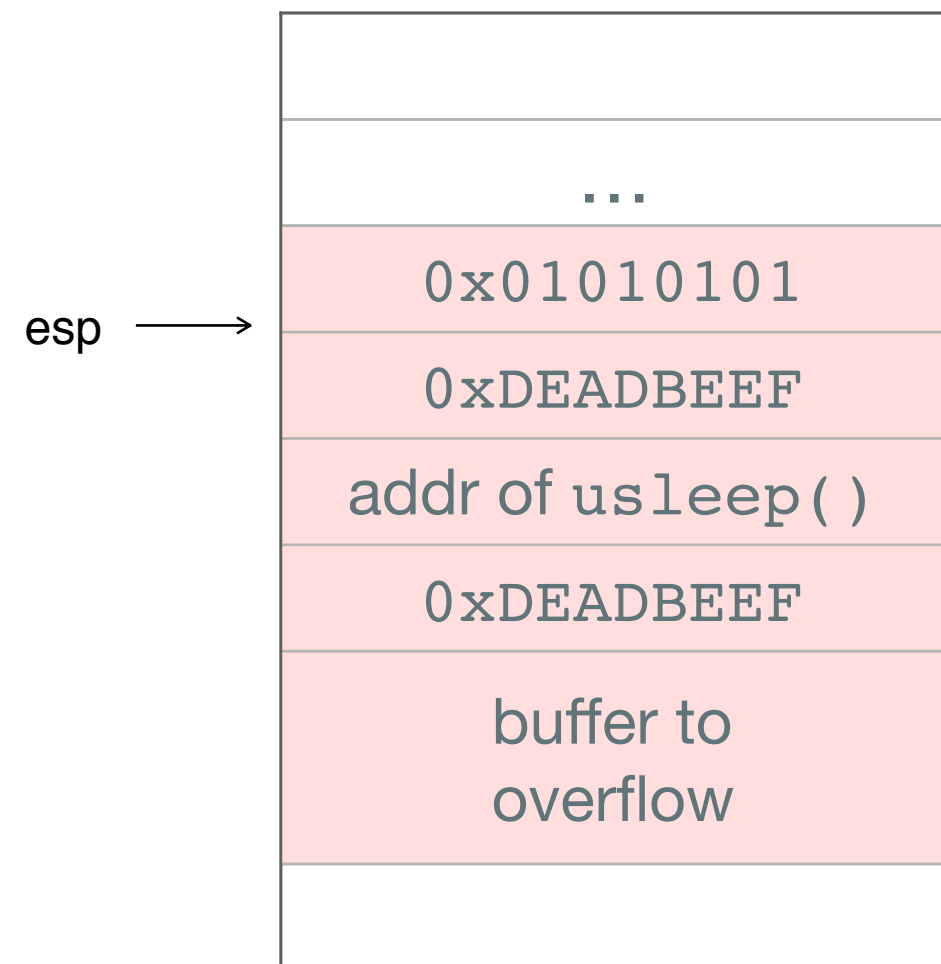
```
pop ebp
```

```
ret
```

```
usleep:
```

```
...
```

```
eip → ret
```



Finding base of mapped region

```
ap_getline:
```

```
...
```

```
pop ebp
```

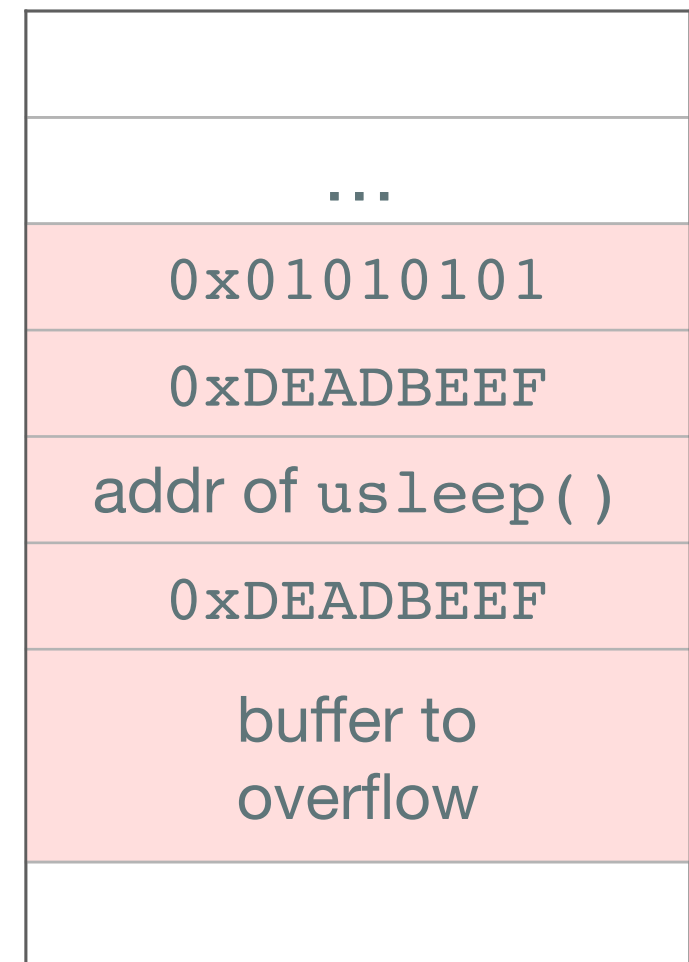
```
ret
```

```
usleep:
```

```
...
```

```
ret
```

esp →



eip → 0xDEADBEEF

Finding base of mapped region

```
ap_getline:
```

```
...
```

```
pop ebp
```

```
ret
```

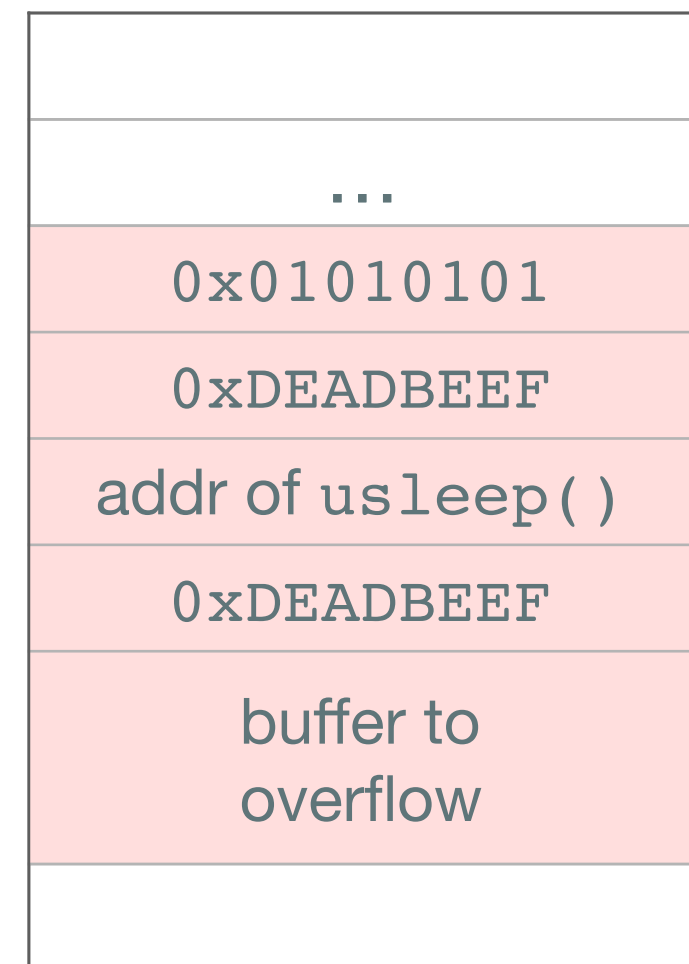
```
usleep:
```

```
...
```

```
ret
```

SEGFault!

esp →



eip → 0xDEADBEEF

Derandomizing ASLR

- What is the success probability?
 -
- Do we need to derandomize the stack base?
 - A: yes, **B: no**
- Attack works even with PaX ASLR and DEP

Derandomizing ASLR

- **Stage 1: Find base of mapped region**

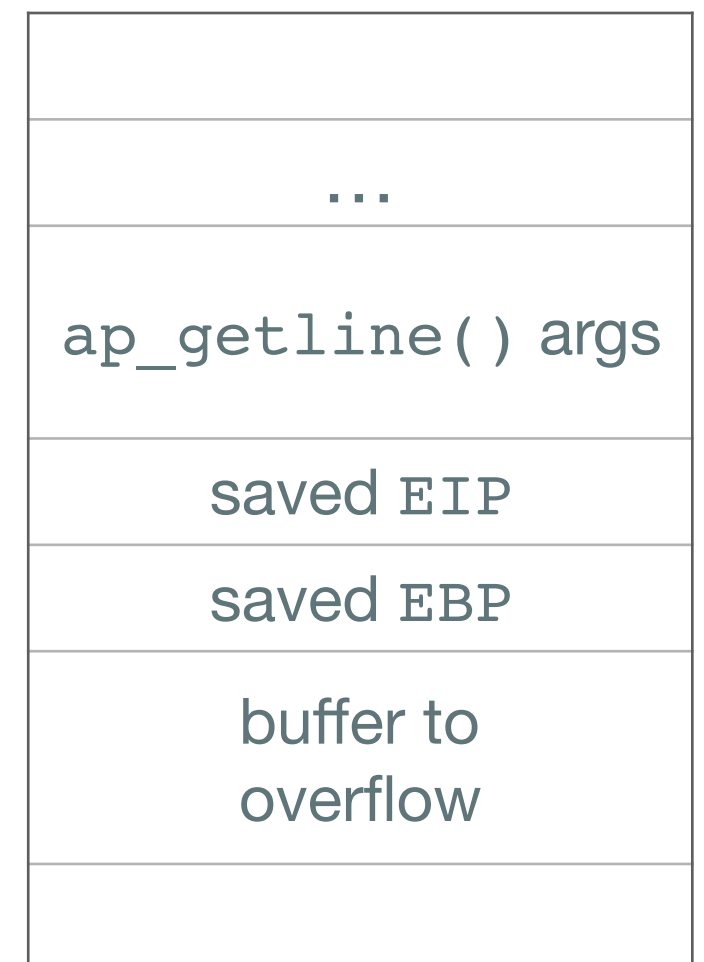
Mapped area:



- **Stage 2: Call system() with command string**

Stage 2

- Overflow buffer in `ap_getline()`
- Pointer to buffer is a local in `ap_getline`
 - Overwrite saved EIP with address of (any) `ret` instruction in `libc`
 - Repeat until address of attack command string on the stack
- Append address of `system()`



Stage 2

top of stack (higher addresses)
⋮
ap_getline() arguments
saved EIP
saved EBP
64 byte buffer
⋮
bottom of stack (lower addresses)

Stage 2

top of stack (higher addresses)
⋮
ap_getline() arguments
saved EIP
saved EBP
64 byte buffer
⋮
bottom of stack (lower addresses)



Stage 2

top of stack (higher addresses)
⋮
<code>ap_getline()</code> arguments
saved EIP
saved EBP
64 byte buffer
⋮
bottom of stack (lower addresses)



top of stack (higher addresses)
⋮
pointer into 64 byte buffer
0xdeadbeef
address of <code>system()</code>
address of <code>ret</code> instruction
⋮
address of <code>ret</code> instruction
0xdeadbeef
64 byte buffer (contains shell commands)
⋮
bottom of stack (lower addresses)

Stage 2

already on stack, adjust
esp (w/ rets) to make it
look like arg to system()

top of stack (higher addresses)
⋮
ap_getline() arguments
saved EIP
saved EBP
64 byte buffer
⋮
bottom of stack (lower addresses)



top of stack (higher addresses)
⋮
pointer into 64 byte buffer
0xdeadbeef
address of system()
address of ret instruction
⋮
address of ret instruction
0xdeadbeef
64 byte buffer (contains shell commands)
⋮
bottom of stack (lower addresses)

Summary: return to libc

- If stack not executable, what can we do?
 - Use existing program code! return-to-libc
- Search executable for code
 - E.g. if executable calls `exec("/bin/sh")`, jump there
- Need known executable
 - Usually not a problem, can work around this

**Employees must
wash hands before
returning to libc**



What if there is no code that does what we want?

The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)

Hovav Shacham*
hovav@cs.ucsd.edu

Return-Oriented Programming

- Idea: make shellcode out of existing code
- Gadgets: code sequences ending in ret instruction
 - Overwrite saved EIP on stack to pointer to first gadget, then second gadget, etc.
- Where do you often find ret instructions?
 -
 -

Return-Oriented Programming

- Idea: make shellcode out of existing code
- Gadgets: code sequences ending in ret instruction
 - Overwrite saved EIP on stack to pointer to first gadget, then second gadget, etc.
- Where do you often find ret instructions?
 - End of function (inserted by compiler)
 -

Return-Oriented Programming

- Idea: make shellcode out of existing code
- Gadgets: code sequences ending in ret instruction
 - Overwrite saved EIP on stack to pointer to first gadget, then second gadget, etc.
- Where do you often find ret instructions?
 - End of function (inserted by compiler)
 - Any sequence of executable memory ending in 0xc3

x86 instructions

- Variable length!
- Can begin on any byte boundary!

So?


- b8 01 00 00 00 5b c9 c3
 - mov eax,0x1 → pop ebx → leave → ret
- 00 00 5b c9 c3
 - add BYTE PTR [eax],al → pop ebx → leave → ret
- 00 5b c9 c3
 - add BYTE PTR [eax-0x37],bl → ret

compy% otool -t /bin/ls
/bin/ls:
(__TEXT,__text) section

00000000100001478	6a	00	48	89	e5	48	83	e4	f0	48	8b	7d	08	48	8d	75
00000000100001488	10	89	fa	83	c2	01	c1	e2	03	48	01	f2	48	89	d1	eb
00000000100001498	04	48	83	c1	08	48	83	39	00	75	f6	48	83	c1	08	e8
000000001000014a8	58	0f	00	00	89	c7	e8	1b	39	00	00	f4	55	48	89	e5
000000001000014b8	48	8d	47	68	48	8d	7e	68	48	89	c6	c9	e9	01	3a	00
000000001000014c8	00	55	48	89	e5	48	83	c6	68	48	83	c7	68	c9	e9	ef
000000001000014d8	39	00	00	55	48	89	e5	53	48	89	f1	48	8b	56	60	48
000000001000014e8	8b	47	60	48	8b	58	30	48	39	5a	30	7f	1d	7c	22	48
000000001000014f8	8b	58	38	48	39	5a	38	7f	11	7c	16	48	8d	77	68	48
00000000100001508	8d	79	68	5b	c9	e9	b8	39	00	00	b8	ff	ff	ff	ff	eb
00000000100001518	05	b8	01	00	00	00	5b	c9	c3	55	48	89	e5	48	8b	56
00000000100001528	60	48	8b	47	60	48	8b	48	50	48	39	4a	50	7f	1c	7c
00000000100001538	21	48	8b	48	58	48	39	4a	58	7f	10	7c	15	48	83	c6
00000000100001548	68	48	83	c7	68	c9	e9	77	39	00	00	b8	01	00	00	00
00000000100001558	eb	05	b8	ff	ff	ff	ff	c9	c3	55	48	89	e5	53	48	8b
00000000100001568	56	60	48	8b	47	60	b9	01	00	00	00	48	8b	58	60	48
00000000100001578	39	5a	60	7f	18	7d	07	b9	ff	ff	ff	ff	eb	0f	48	83
00000000100001588	c6	68	48	83	c7	68	5b	c9	e9	35	39	00	00	89	c8	5b
00000000100001598	c9	c3	55	48	89	e5	48	8b	56	60	48	8b	47	60	48	8b
000000001000015a8	48	40	48	39	4a	40	7f	1c	7c	21	48	8b	48	48	48	39
000000001000015b8	4a	48	7f	10	7c	15	48	83	c6	68	48	83	c7	68	c9	e9
000000001000015c8	fe	38	00	00	b8	01	00	00	00	eb	05	b8	ff	ff	ff	ff

What does a gadget look like?


- Gadget for loading a constant
 - Arrange the constant to load to be just past the return address
 - Return to gadget that pops a value and returns.



What does this gadget do?

eip \longrightarrow a_1 : pop eax;
 a_2 : ret
 a_3 : pop ebx;
 a_4 : ret

Code




eax					
ebx					
eip	a_1				

time \longrightarrow

What does this gadget do?

eip \longrightarrow a_1 : pop eax;
 a_2 : ret
 a_3 : pop ebx;
 a_4 : ret

Code



Stack


eax		v_1			
ebx					
eip	a_1	a_2			

time \longrightarrow

What does this gadget do?

eip →
a₁: pop eax;
a₂: ret
a₃: pop ebx;
a₄: ret

Code

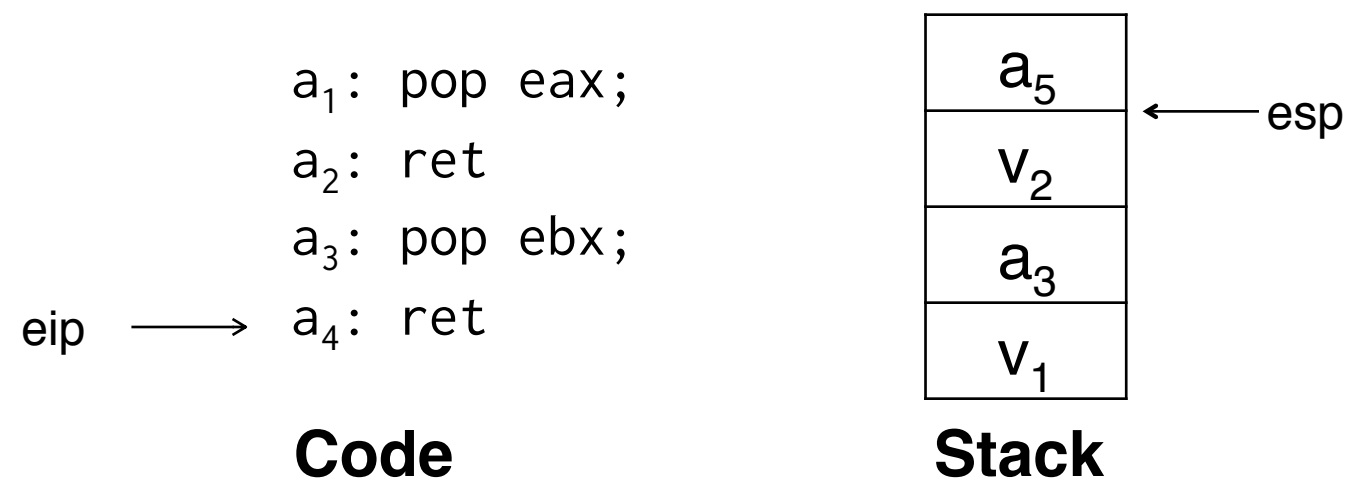


Stack

eax		v ₁	v ₁		
ebx					
eip	a ₁	a ₂	a ₃		

time →

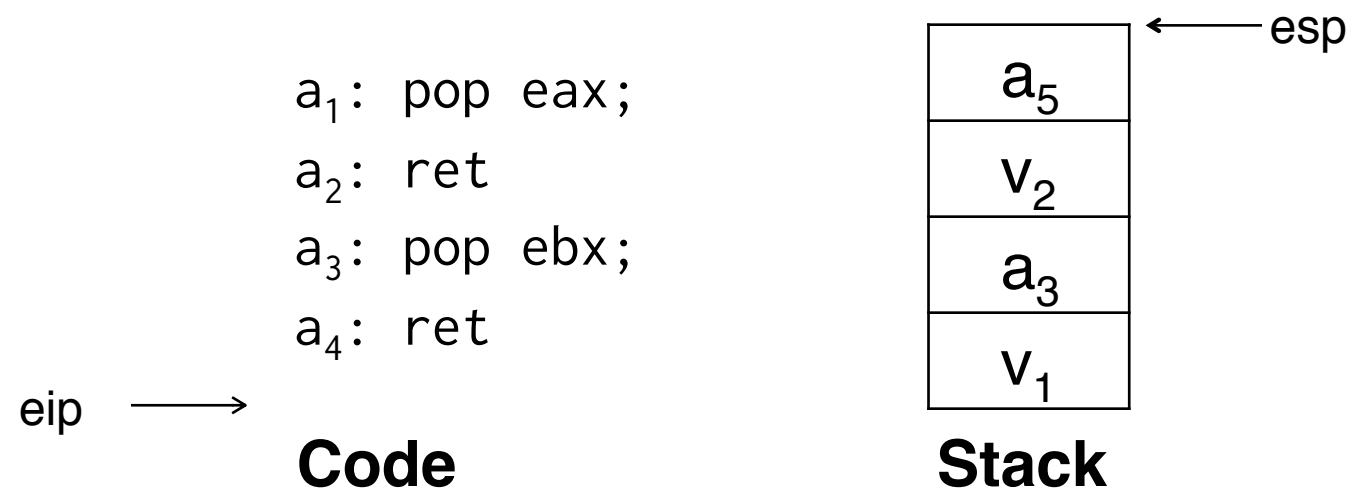
What does this gadget do?



eax		v ₁	v ₁	v ₁	
ebx				v ₂	
eip	a ₁	a ₂	a ₃	a ₄	

time →

What does this gadget do?



eax		v ₁	v ₁	v ₁	v ₁
ebx				v ₂	v ₂
eip	a ₁	a ₂	a ₃	a ₄	a ₅

time →




Figure 5: Simple add into %eax.




Figure 10: An infinite loop by means of an unconditional jump.

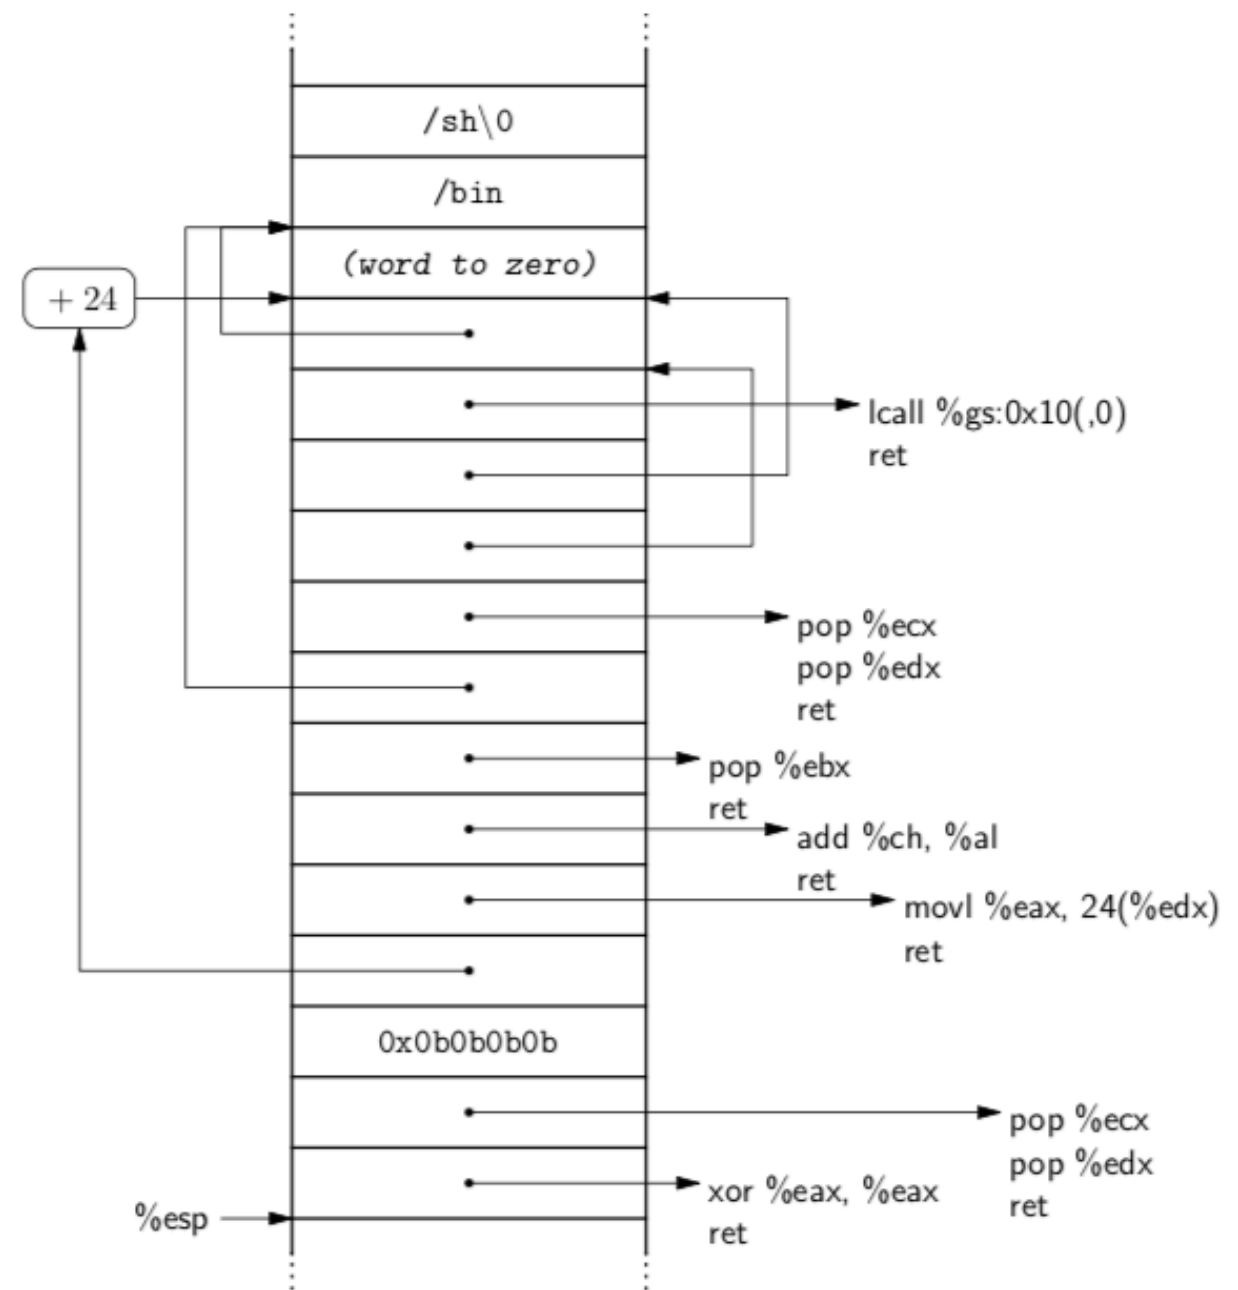


Figure 16: Shellcode.

	“Normal”	Return-oriented
Instruction pointer	eip	
No-op	nop	
Unconditional jump	jmp address	
Conditional jump	jnz address	set esp to address of gadget if some condition is met; ret
Variables	memory and registers	mostly memory
Inter-instruction (inter-gadget) register and memory interaction	minimal, mostly explicit; e.g., adding two registers only affects the destination register	can be complex; e.g., adding two registers may involve modifying many registers which impacts other gadgets

	“Normal”	Return-oriented
Instruction pointer	eip	esp
No-op	nop	
Unconditional jump	jmp address	
Conditional jump	jnz address	set esp to address of gadget if some condition is met; ret
Variables	memory and registers	mostly memory
Inter-instruction (inter-gadget) register and memory interaction	minimal, mostly explicit; e.g., adding two registers only affects the destination register	can be complex; e.g., adding two registers may involve modifying many registers which impacts other gadgets

	“Normal”	Return-oriented
Instruction pointer	eip	esp
No-op	nop	ret
Unconditional jump	jmp address	
Conditional jump	jnz address	set esp to address of gadget if some condition is met; ret
Variables	memory and registers	mostly memory
Inter-instruction (inter-gadget) register and memory interaction	minimal, mostly explicit; e.g., adding two registers only affects the destination register	can be complex; e.g., adding two registers may involve modifying many registers which impacts other gadgets

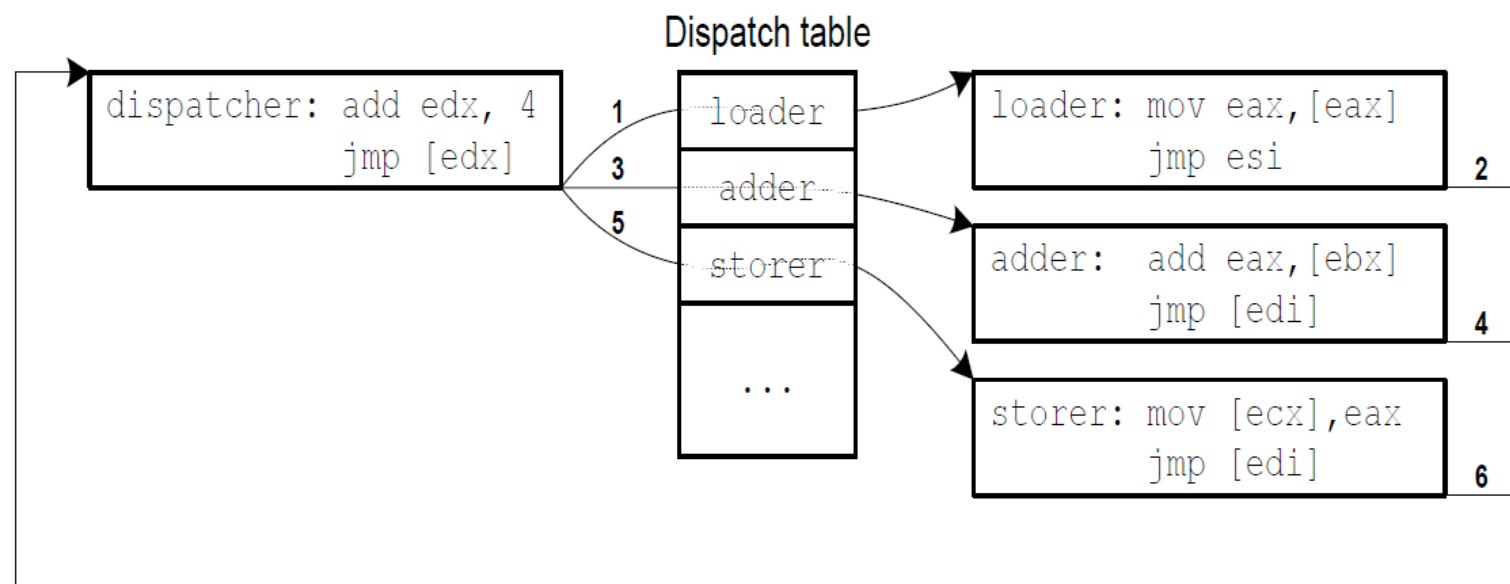
	“Normal”	Return-oriented
Instruction pointer	eip	esp
No-op	nop	ret
Unconditional jump	jmp address	set esp to addr of gadget; ret
Conditional jump	jnz address	set esp to address of gadget if some condition is met; ret
Variables	memory and registers	mostly memory
Inter-instruction (inter- gadget) register and memory interaction	minimal, mostly explicit; e.g., adding two registers only affects the destination register	can be complex; e.g., adding two registers may involve modifying many registers which impacts other gadgets

Return-Oriented Programming

not even really about “returns”...

Jump-Oriented Programming

- Identify gadgets ending in indirect jumps.
- Use a “dispatcher gadget” to combine them.
- Dispatch table used in place of stack





Hacking Blind

Andrea Bittau, Adam Belay, Ali Mashtizadeh, David Mazières, Dan Boneh

Stanford University