



CSE 127: Computer Security

Stack Buffer Overflows

Deian Stefan

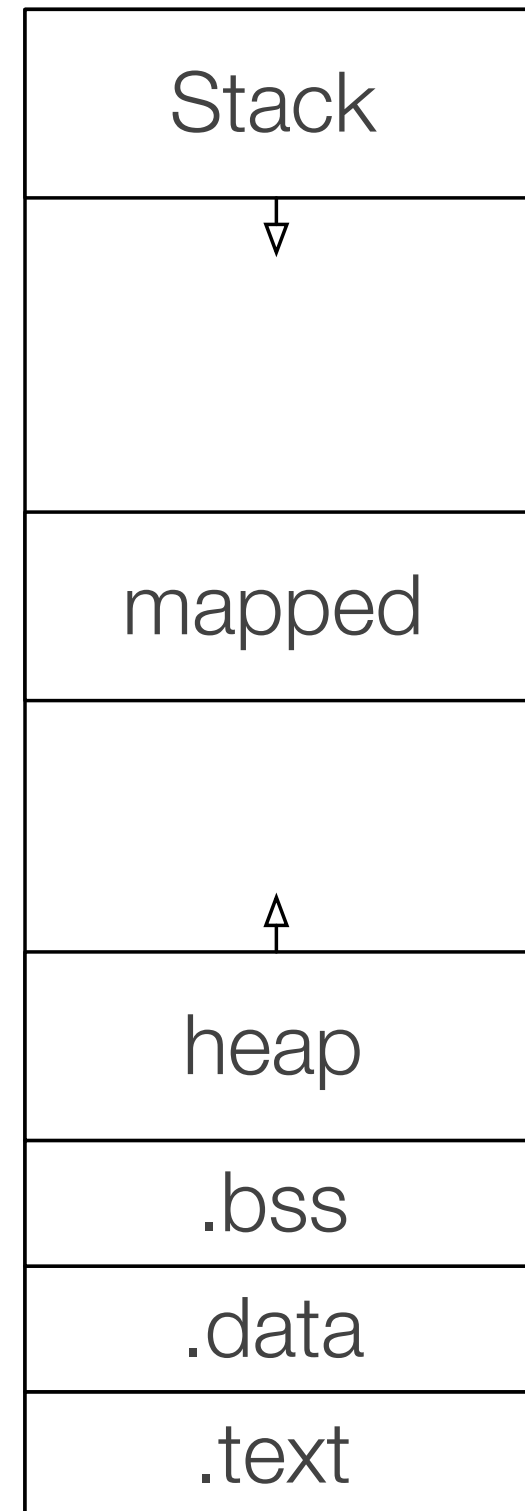
Slides adopted from Kirill Levchenko

Today

- Review: stack overflow attacks
- Shellcode
- Defenses

Process Memory Layout

- Stack
- Heap
- Data
 - Static variables
- Text
 - Executable code



The Stack

- Function local variables
- Function arguments
- Control state

The Stack

- Stack divided into **frames**
 - Frame stores locals and args to called functions
- **Stack pointer** points to
 - x86: Stack grows down (from high to low addresses)
 - x86: Stored in **ESP** register
- **Frame pointer** points to
 - Also called **base pointer**
 - x86: Stored in **EBP** register

The Stack

- Stack divided into **frames**
 - Frame stores locals and args to called functions
- **Stack pointer** points to top of stack
 - x86: Stack grows down (from high to low addresses)
 - x86: Stored in **ESP** register
- **Frame pointer** points to caller's stack frame
 - Also called **base pointer**
 - x86: Stored in **EBP** register

Function Call Example

```
int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}
```

_foobar:

```
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
push    ebp
; From now on, ebp points to the current stack
; frame of the function
mov     ebp, esp
; Make space on the stack for local variables
sub     esp, 16
; eax <-- a. eax += 2. then store eax in x
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax
; eax <-- b. eax += 3. then store eax in yy
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax
; eax <-- c. eax += 4. then store eax in zz
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax
; add xx + yy + zz and store it in sum
mov     eax, DWORD PTR [ebp-8]
mov     edx, DWORD PTR [ebp-4]
lea     eax, [edx+eax]
add     eax, DWORD PTR [ebp-12]
mov     DWORD PTR [ebp-16], eax
; Compute final result into eax
mov     eax, DWORD PTR [ebp-4]
imul    eax, DWORD PTR [ebp-8]
imul    eax, DWORD PTR [ebp-12]
add     eax, DWORD PTR [ebp-16]
; The leave instruction here is equivalent to:
;   mov esp, ebp ; pop ebp
leave
ret
```



Function Call Example

```
int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}
```

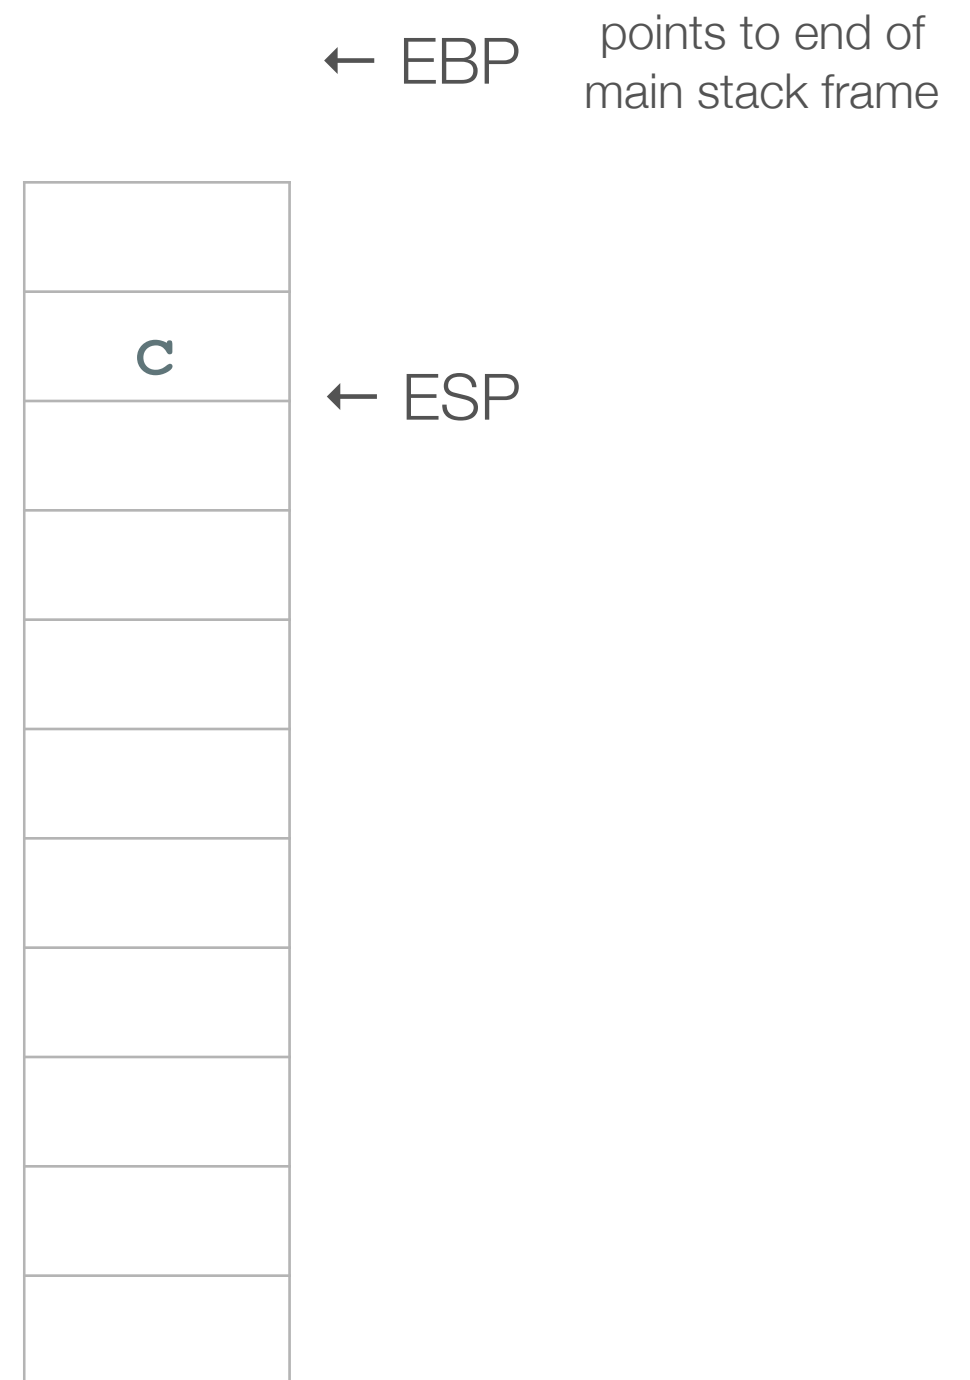
```

int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}

```



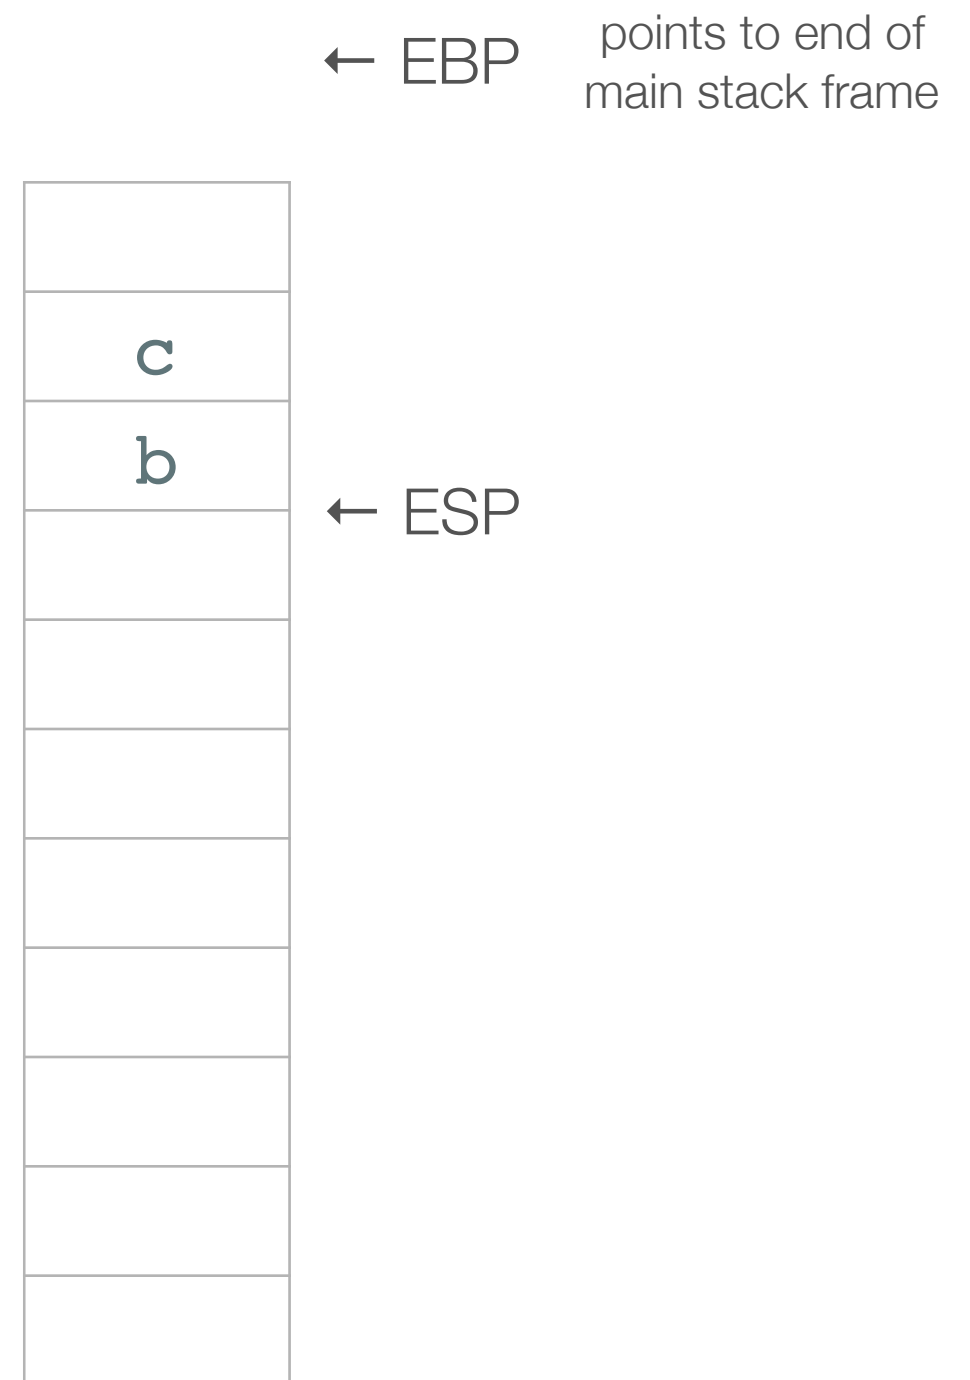
```

int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}

```



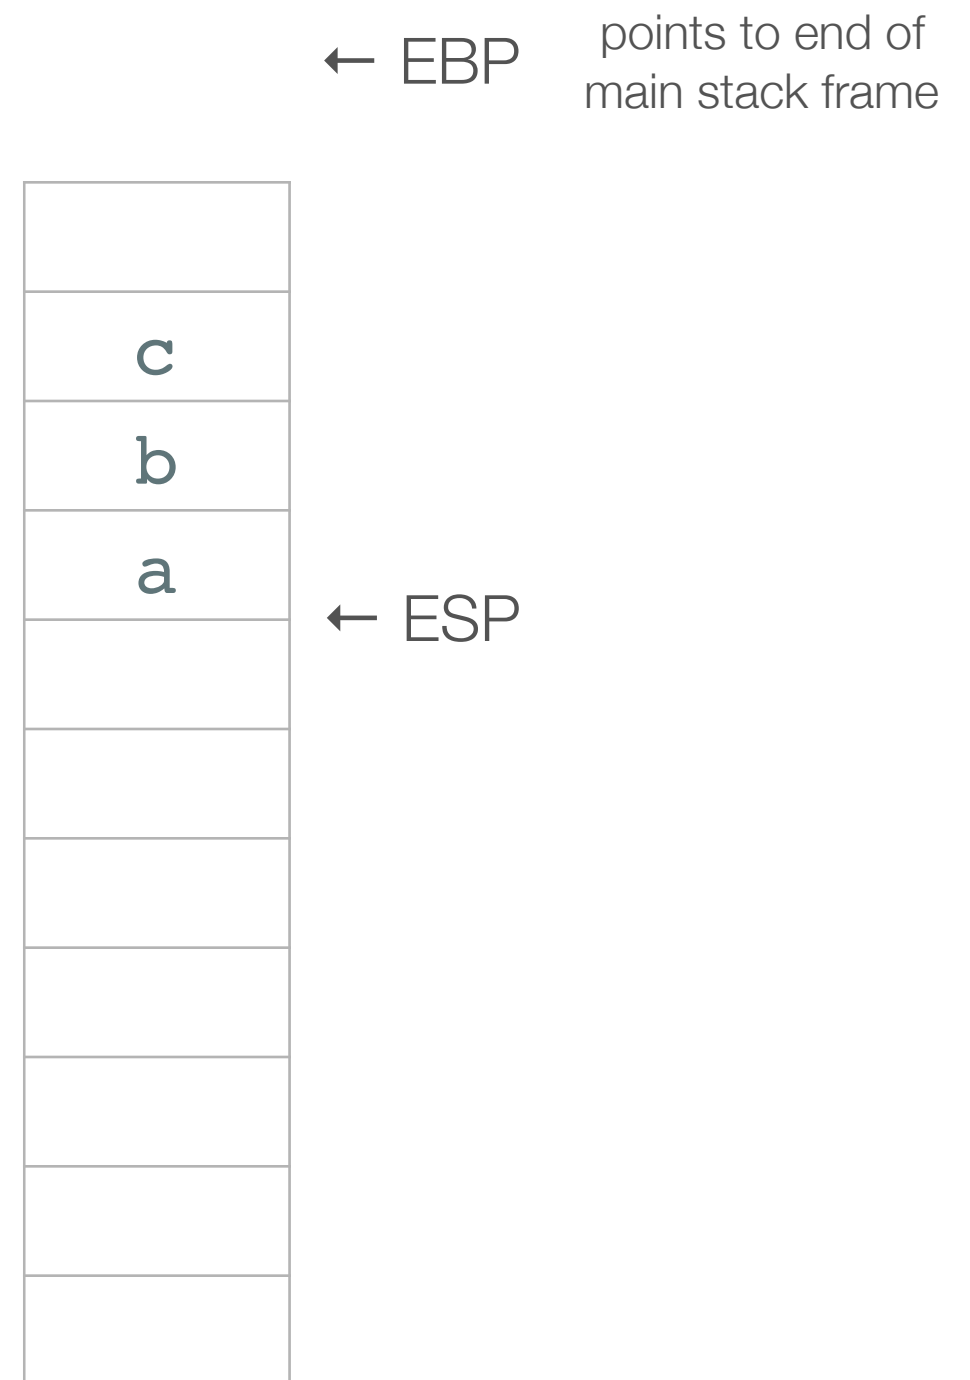
```

int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}

```



```
int foobar(int a, int b, int c)
{
    int xx = a + 2;
    int yy = b + 3;
    int zz = c + 4;
    int sum = xx + yy + zz;

    return xx * yy * zz + sum;
}

int main()
{
    return foobar(77, 88, 99);
}
```



← EBP points to end of main stack frame

← ESP

```

foobar:
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
;
push    ebp

; From now on, ebp points to the current stack
; frame of the function
;
mov     ebp, esp

; Make space on the stack for local variables
;
sub     esp, 16

; eax <-- a. eax += 2. then store eax in xx
;
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax

; eax <-- b. eax += 3. then store eax in yy
;
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

```



← EBP points to end of
main stack frame

← ESP

```

foobar:
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
;
push    ebp

; From now on, ebp points to the current stack
; frame of the function
;
mov     ebp, esp

; Make space on the stack for local variables
;
sub     esp, 16

; eax <-- a. eax += 2. then store eax in xx
;
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax

; eax <-- b. eax += 3. then store eax in yy
;
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

```



← EBP points to end of
main stack frame

← ESP

```

foobar:
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
;
push    ebp

; From now on, ebp points to the current stack
; frame of the function
;
mov     ebp, esp


; Make space on the stack for local variables
;
sub     esp, 16

; eax <-- a. eax += 2. then store eax in xx
;
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax

; eax <-- b. eax += 3. then store eax in yy
;
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

```



← ESP, EBP


```

foobar:
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
;
push    ebp

; From now on, ebp points to the current stack
; frame of the function
;
mov     ebp, esp

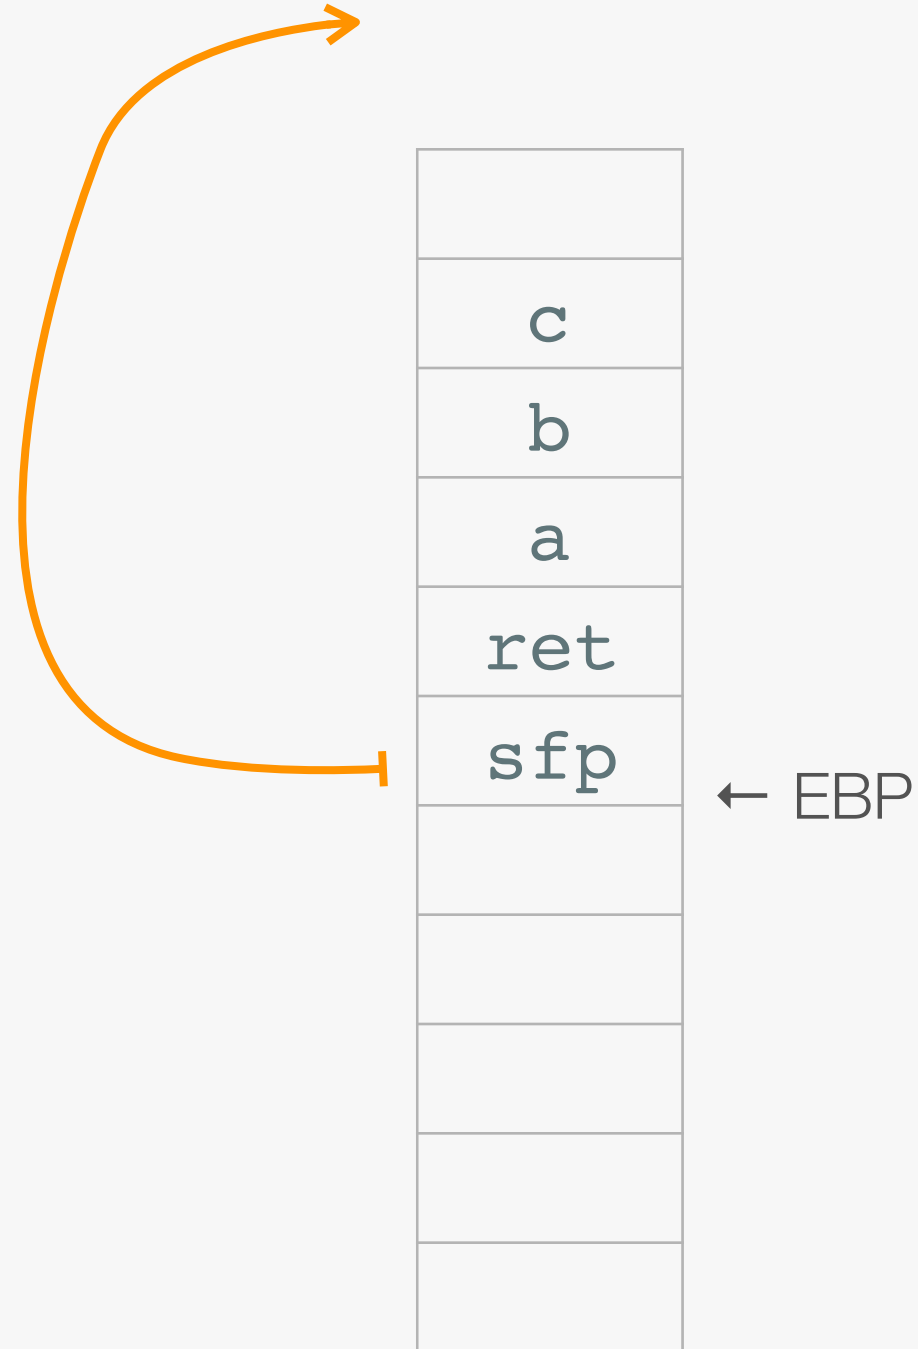
; Make space on the stack for local variables
;
sub     esp, 16

; eax <-- a. eax += 2. then store eax in xx
;
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax

; eax <-- b. eax += 3. then store eax in yy
;
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

```



```

foobar:
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
;
push    ebp

; From now on, ebp points to the current stack
; frame of the function
;
mov     ebp, esp

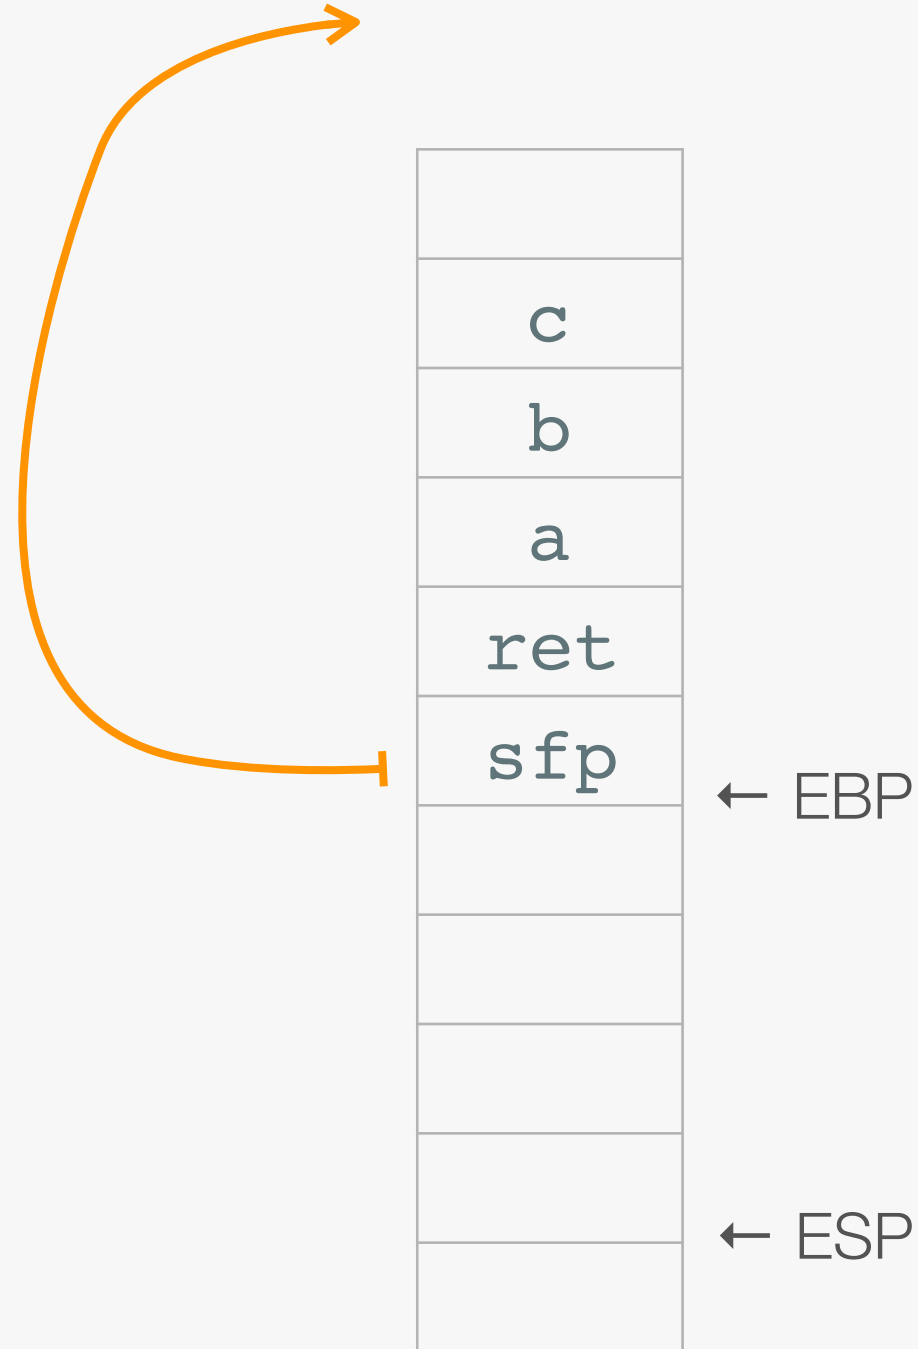
; Make space on the stack for local variables
;
sub     esp, 16

; eax <-- a. eax += 2. then store eax in xx
;
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax

; eax <-- b. eax += 3. then store eax in yy
;
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

```



```

foobar:
; ebp must be preserved across calls. Since
; this function modifies it, it must be
; saved.
;
push    ebp

; From now on, ebp points to the current stack
; frame of the function
;
mov     ebp, esp

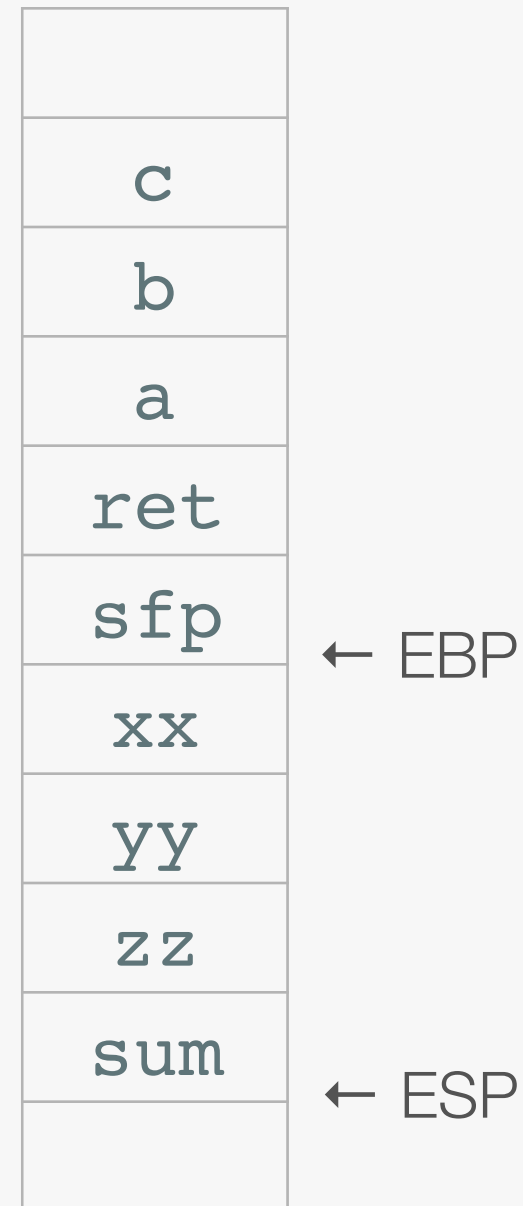
; Make space on the stack for local variables
;
sub     esp, 16

; eax <-- a. eax += 2. then store eax in xx
;
mov     eax, DWORD PTR [ebp+8]
add     eax, 2
mov     DWORD PTR [ebp-4], eax

; eax <-- b. eax += 3. then store eax in yy
;
mov     eax, DWORD PTR [ebp+12]
add     eax, 3
mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

```



```

mov     DWORD PTR [ebp-8], eax

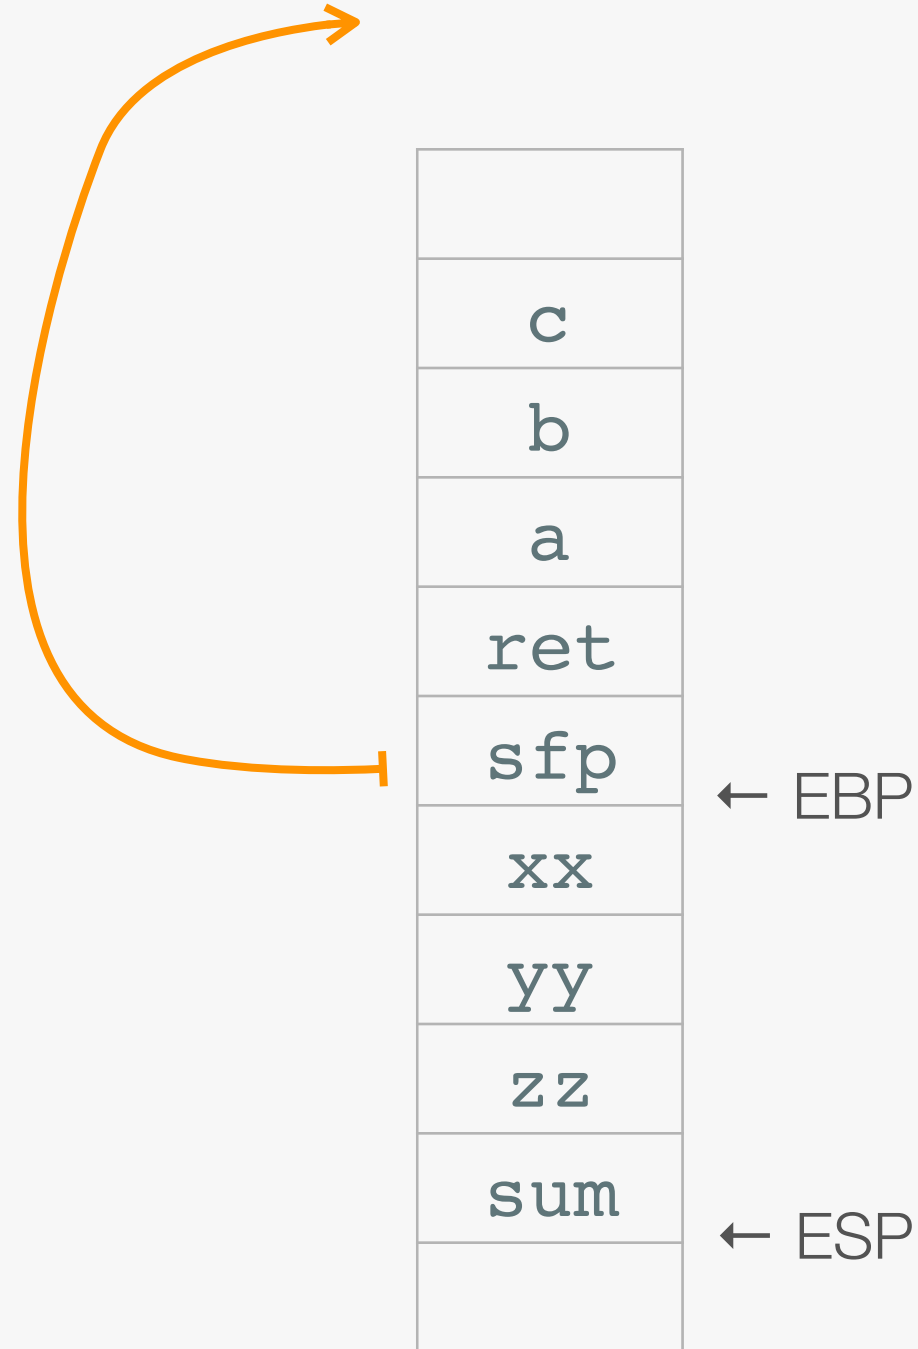
; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

; add xx + yy + zz and store it in sum
;
mov     eax, DWORD PTR [ebp-8]
mov     edx, DWORD PTR [ebp-4]
lea     eax, [edx+eax]
add     eax, DWORD PTR [ebp-12]
mov     DWORD PTR [ebp-16], eax

; Compute final result into eax, which
; stays there until return
;
mov     eax, DWORD PTR [ebp-4]
imul    eax, DWORD PTR [ebp-8]
imul    eax, DWORD PTR [ebp-12]
add     eax, DWORD PTR [ebp-16]

; The leave instruction here is equivalent to:
;
;   mov esp, ebp
;   pop ebp
;
; Which cleans the allocated locals and restores
; ebp.
;
leave
ret

```



```

mov     DWORD PTR [ebp-8], eax

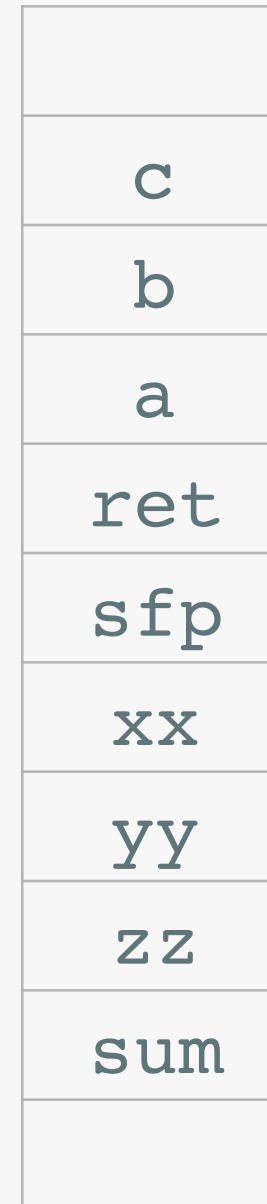
; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

; add xx + yy + zz and store it in sum
;
mov     eax, DWORD PTR [ebp-8]
mov     edx, DWORD PTR [ebp-4]
lea     eax, [edx+eax]
add     eax, DWORD PTR [ebp-12]
mov     DWORD PTR [ebp-16], eax

; Compute final result into eax, which
; stays there until return
;
mov     eax, DWORD PTR [ebp-4]
imul    eax, DWORD PTR [ebp-8]
imul    eax, DWORD PTR [ebp-12]
add     eax, DWORD PTR [ebp-16]

; The leave instruction here is equivalent to:
;
;   mov esp, ebp
;   pop ebp
;
; Which cleans the allocated locals and restores
; ebp.
;
leave
ret

```



← EBP, ESP


```

mov     DWORD PTR [ebp-8], eax

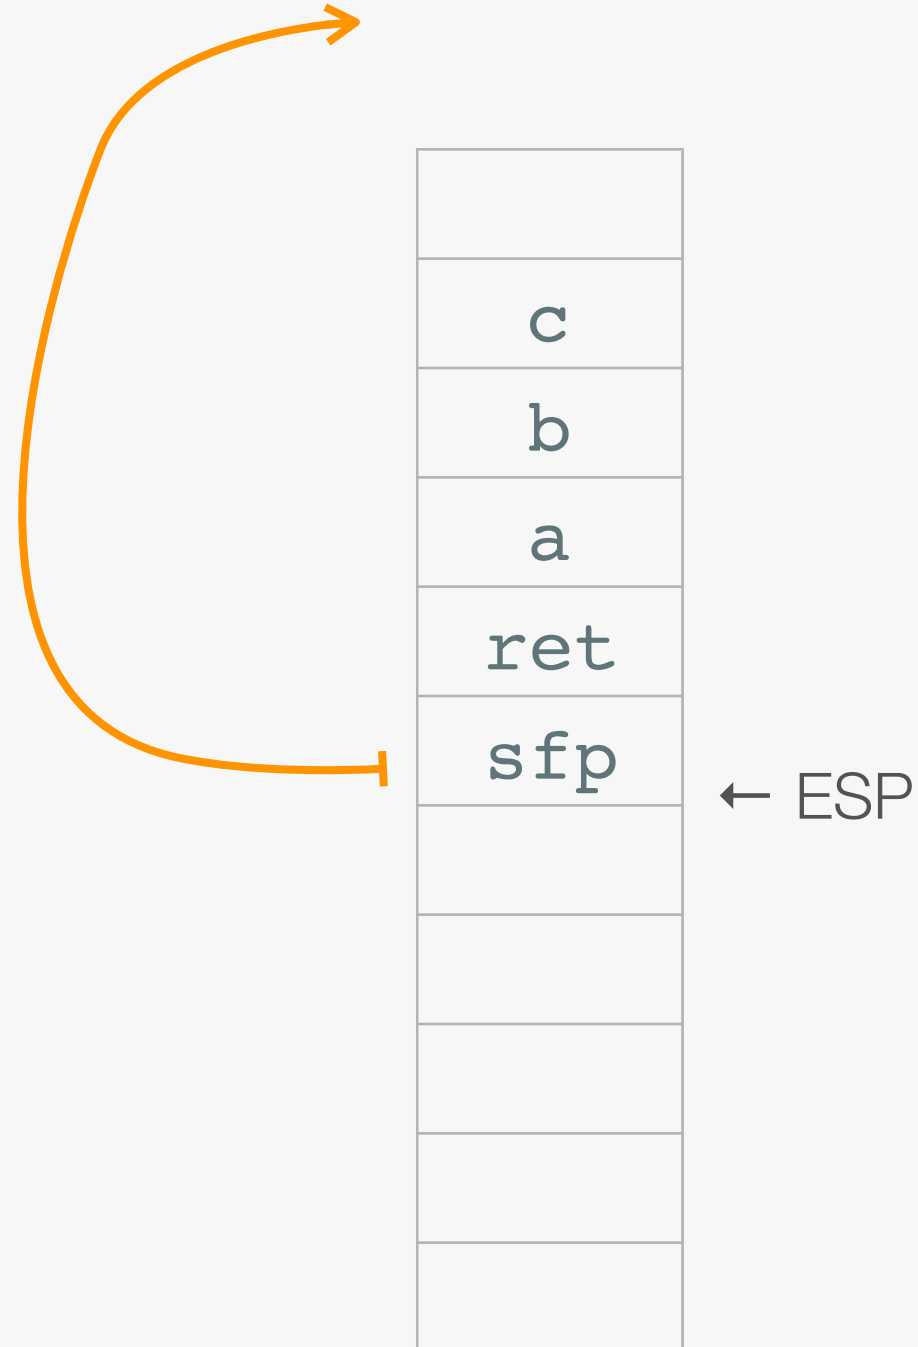
; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

; add xx + yy + zz and store it in sum
;
mov     eax, DWORD PTR [ebp-8]
mov     edx, DWORD PTR [ebp-4]
lea     eax, [edx+eax]
add     eax, DWORD PTR [ebp-12]
mov     DWORD PTR [ebp-16], eax

; Compute final result into eax, which
; stays there until return
;
mov     eax, DWORD PTR [ebp-4]
imul    eax, DWORD PTR [ebp-8]
imul    eax, DWORD PTR [ebp-12]
add     eax, DWORD PTR [ebp-16]

; The leave instruction here is equivalent to:
;
;   mov esp, ebp
;   pop ebp
;
; Which cleans the allocated locals and restores
; ebp.
;
leave
ret

```



```

mov     DWORD PTR [ebp-8], eax

; eax <-- c. eax += 4. then store eax in zz
;
mov     eax, DWORD PTR [ebp+16]
add     eax, 4
mov     DWORD PTR [ebp-12], eax

; add xx + yy + zz and store it in sum
;
mov     eax, DWORD PTR [ebp-8]
mov     edx, DWORD PTR [ebp-4]
lea     eax, [edx+eax]
add     eax, DWORD PTR [ebp-12]
mov     DWORD PTR [ebp-16], eax

; Compute final result into eax, which
; stays there until return
;
mov     eax, DWORD PTR [ebp-4]
imul    eax, DWORD PTR [ebp-8]
imul    eax, DWORD PTR [ebp-12]
add     eax, DWORD PTR [ebp-16]

; The leave instruction here is equivalent to:
;
;   mov esp, ebp
;   pop ebp
;
; Which cleans the allocated locals and restores
; ebp.
;
leave
ret

```




← EBP points to end of main stack frame

← ESP

With Buffers

example1.c:

```
-----  
void function(int a, int b, int c) {  
    char buffer1[5];  
    char buffer2[10];  
}  
  
void main() {  
    function(1,2,3);  
}  
-----
```



With Buffers

```
buffer      sfp  ret  *str  
[           ][     ][     ][     ]
```

- strcpy will copy memory from str to buffer
- So?



example2.c

```
-----  
void function(char *str) {  
    char buffer[16];  
  
    strcpy(buffer, str);  
}  
  
void main() {  
    char large_string[256];  
    int i;  
  
    for( i = 0; i < 255; i++)  
        large_string[i] = 'A';  
  
    function(large_string);  
}  
-----
```

With Buffers

```
buffer      sfp   ret   *str
[           ][     ][     ][     ]
```

- `strcpy` will copy memory from `str` to `buffer` until `'\0'`
- So?
 - if length of string longer than `buffer`, `strcpy` will copy string over `sfp` and `ret`

example2.c

```
-----
void function(char *str) {
    char buffer[16];

    strcpy(buffer, str);
}

void main() {
    char large_string[256];
    int i;

    for( i = 0; i < 255; i++)
        large_string[i] = 'A';

    function(large_string);
}
-----
```

Stack Buffer Overflow

- If source string of `strcpy` controlled by attacker (and destination is on the stack)
 - Attacker gets to control where the function returns by overwriting *ret*
 - Attacker gets to transfer control to anywhere!
- Where do you jump?

Taking Control

- Let's jump to code that does what we want
- **Where?** We have control of string!
 - Put code in string
 - Jump to start of string

bottom of memory	DDDDDDDDDEEEEEEEEEEEEEEE	EEEE	FFFF	FFFF	FFFF	FFFF	top of memory
	89ABCDEF0123456789AB	CDEF	0123	4567	89AB	CDEF	
	buffer	sfp	ret	a	b	c	
<-----	[][][][][][]
top of stack							bottom of stack

Taking Control

- Let's jump to code that does what we want
- **Where?** We have control of string!
 - Put code in string
 - Jump to start of string

bottom of memory	DDDDDDDDDEEEEEEEEEEEEEEE	EEEE	FFFF	FFFF	FFFF	FFFF	top of memory
	89ABCDEF0123456789AB	CDEF	0123	4567	89AB	CDEF	
	buffer	sfp	ret	a	b	c	
<-----	[SSSSSSSSSSSSSSSSSSSSSSSSSSSSSS][SSSS][0xD8][0x01][0x02][0x03]						
top of stack							bottom of stack

Shellcode

- **Shellcode:** small code fragment that receives initial control in an control flow hijack exploit
 - Control flow hijack: taking control of instruction ptr
- Earliest attacks used shellcode to exec a shell
 - Target a setuid root program, gives you root shell

Shellcode

```
void main() {  
    char *name[2];  
  
    name[0] = "/bin/sh";  
    name[1] = NULL;  
    execve(name[0], name, NULL);  
}
```

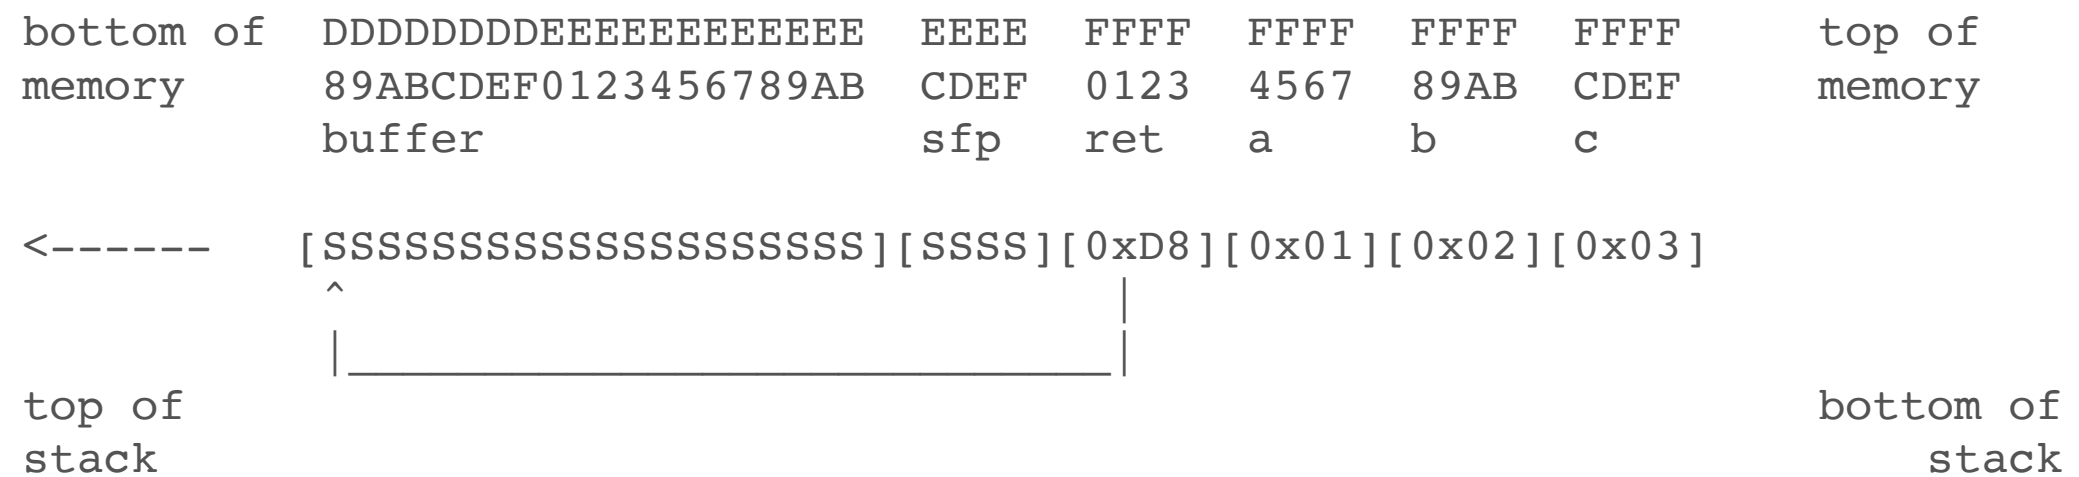
Shellcode

- Can we just take output from gcc/clang?
 - A: yes B: no

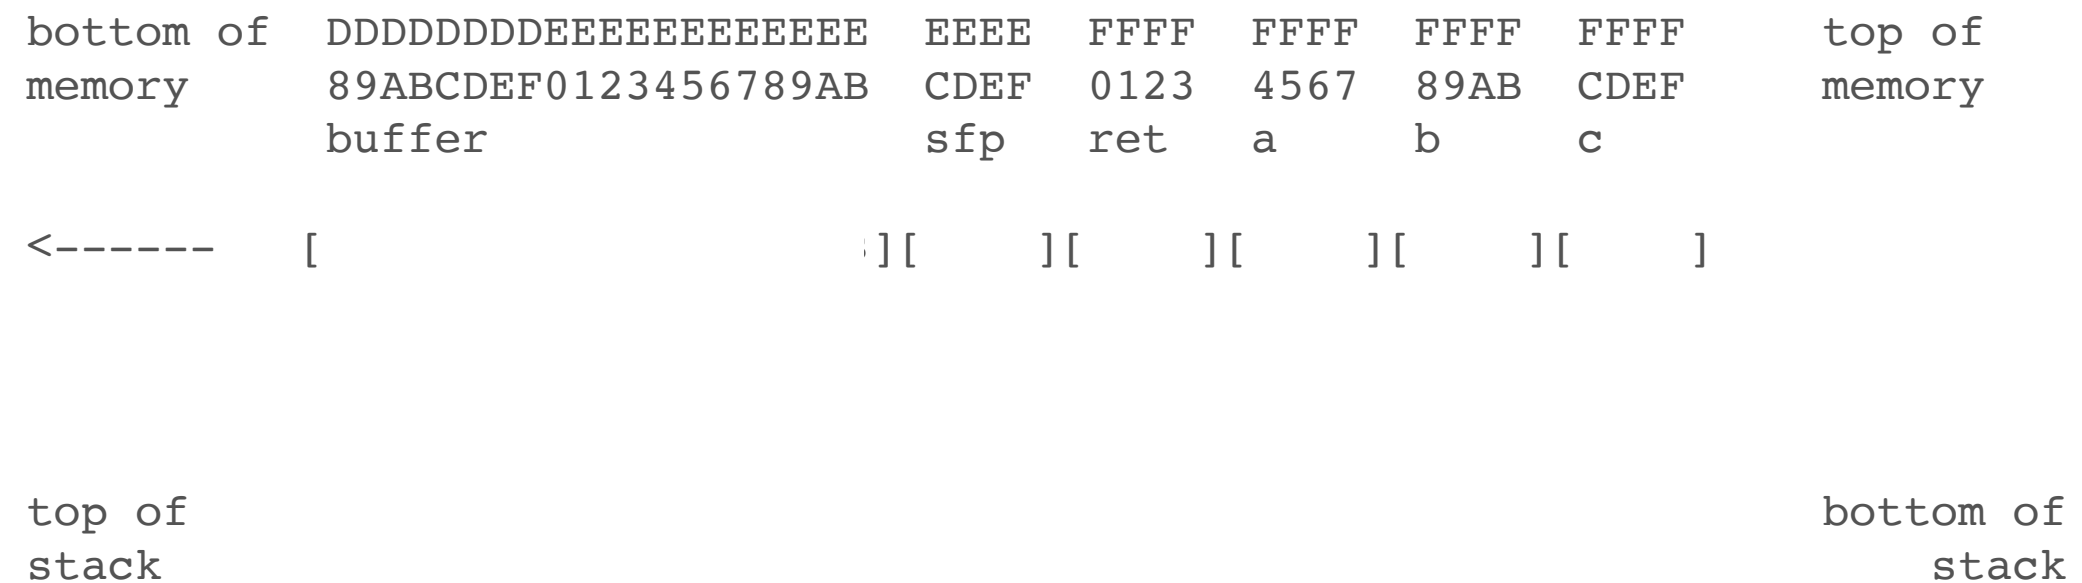
Shellcode

- There some restrictions
 - Shellcode cannot contain null characters '\0'
 - Why not?
 - Recipe:
 - a) Have the null terminated string `"/bin/sh"` somewhere in memory.
 - b) Have the address of the string `"/bin/sh"` somewhere in memory followed by a null long word.
 - c) Copy `0xb` into the `EAX` register.
 - d) Copy the address of the address of the string `"/bin/sh"` into the `EBX` register.
 - e) Copy the address of the string `"/bin/sh"` into the `ECX` register.
 - f) Copy the address of the null long word into the `EDX` register.
 - g) Execute the `int $0x80` instruction.

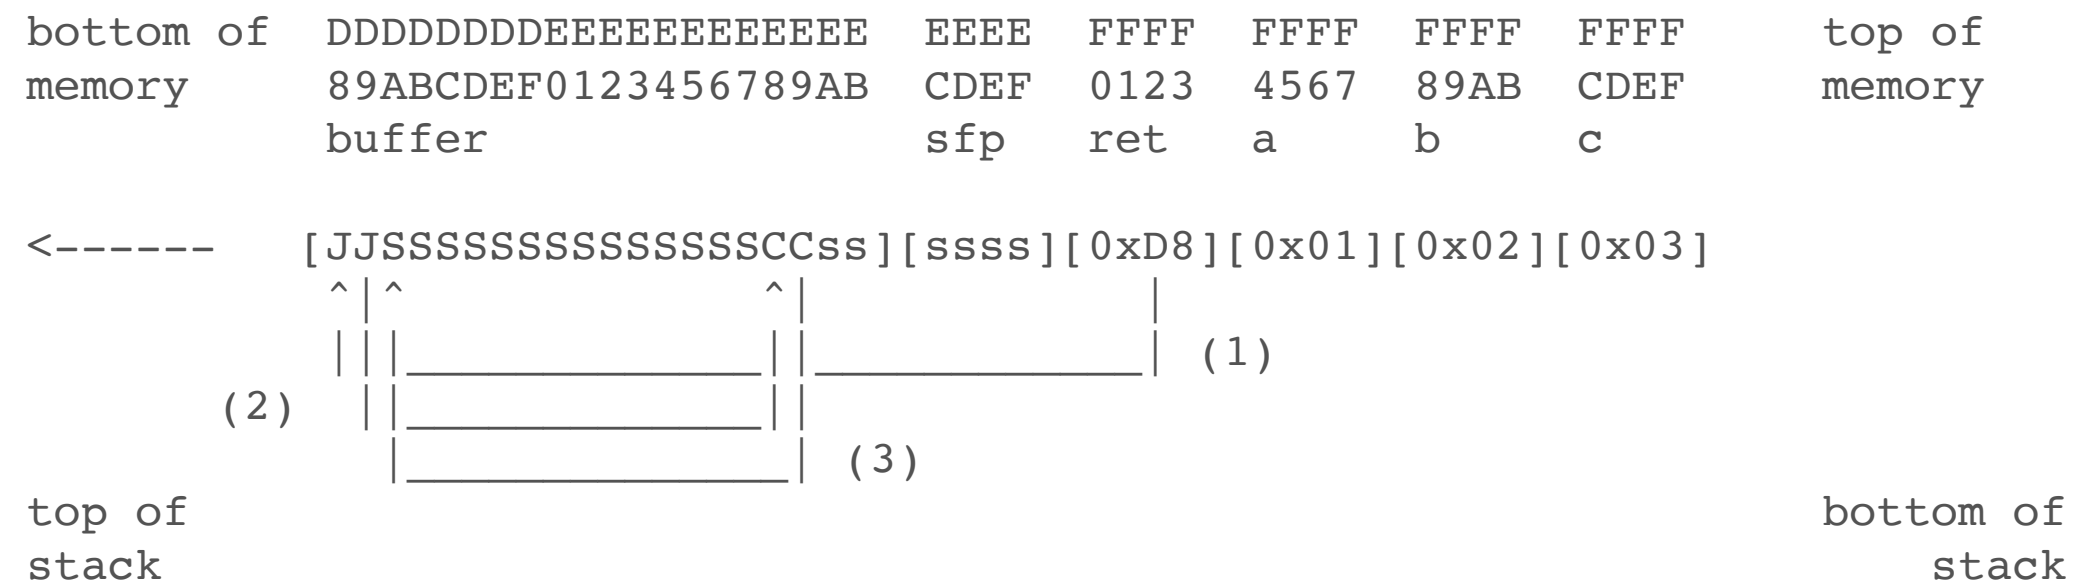
Why does this not really work?



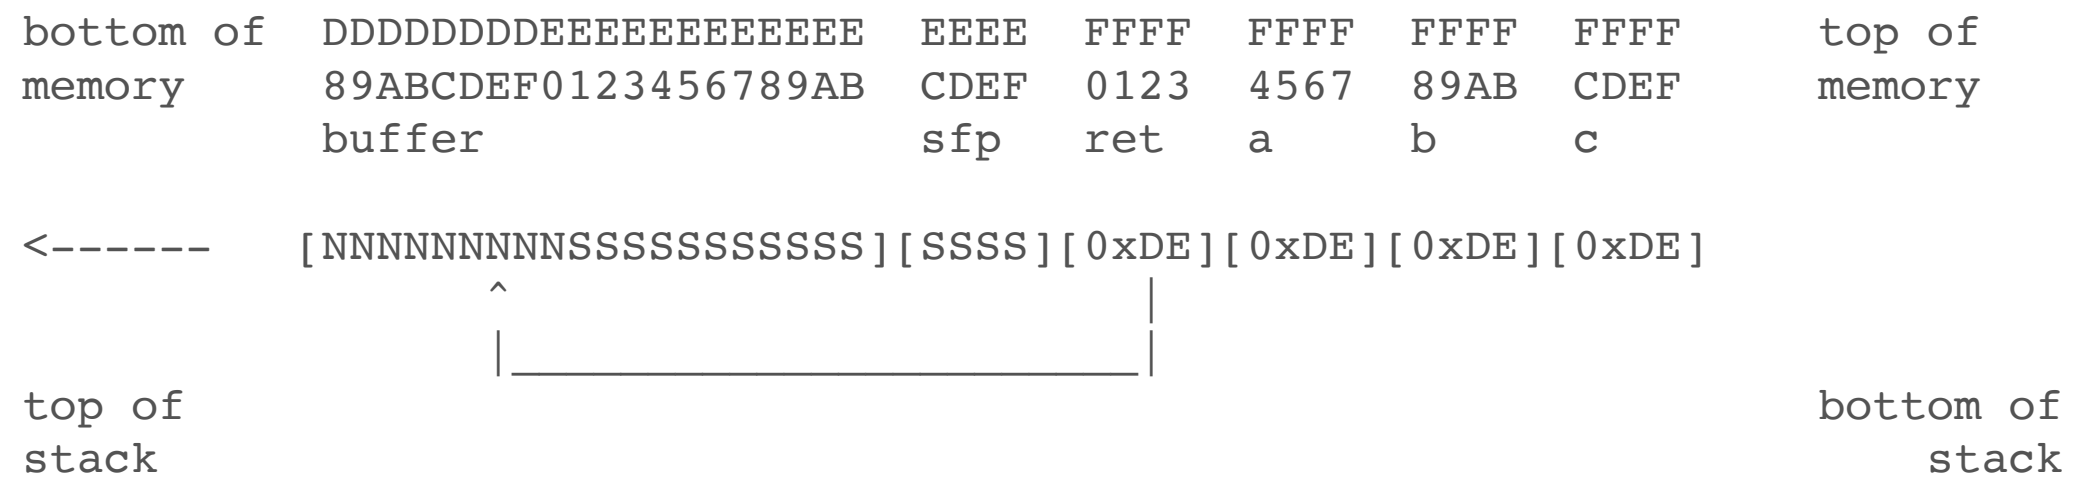
How can we address this?



How can we address this?



How can we address this?



Metasploit to the rescue!

```
msf payload(shell_bind_tcp) > generate -h
```

```
Usage: generate [options]
```

```
Generates a payload.
```

OPTIONS:

- E Force encoding.
- b The list of characters to avoid: '\x00\xff'
- e The name of the encoder module to use.
- f The output file name (otherwise stdout)
- h Help banner.
- i the number of encoding iterations.
- k Keep the template executable functional
- o A comma separated list of options in VAR=VAL format.
- p The Platform for output.
- s NOP sled length.
- t The output format: raw,ruby,rb,perl,pl,c,js_be,js_le,java,dll,exe,exe-small,elf,macho,vba,vml
- x The executable template to use

Buffer Overflow Defenses

- Avoid unsafe functions
- Stack canary
- Separate control stack
- Address Space Layout Randomization (ASLR)
- Memory writable or executable, not both (W^X)
- Control flow integrity (CFI)

Avoiding Unsafe Functions

- strcpy, strcat, gets, etc.
- **Plus:** Good idea in general
- **Minus:** Requires manual code rewrite
- **Minus:** Non-library functions may be vulnerable
 - E.g. user creates their own strcpy
- **Minus:** No guarantee you found everything
- **Minus:** alternatives are also error-prone

Even printf is tricky

If buf is under control of attacker
is: `printf(“%s\n”, buf)` safe?

A: yes, B: no

Even printf is tricky

If buf is under control of attacker
is: `printf(buf)` safe?

A: yes, B: no

Even printf is tricky

Is `printf("%s\n)` safe?

A: yes, B: no

Even printf is tricky

printf can be used to read and write memory

➡ control flow hijacking!

Exploiting Format String Vulnerabilities

scut / team teso

September 1, 2001

Buffer Overflow Defenses

- Avoid unsafe functions
- Stack canary
- Separate control stack
- Address Space Layout Randomization (ASLR)
- Memory writable or executable, not both (W^X)
- Control flow integrity (CFI)

Stack Canary

- Special value placed before return address
 - Secret random value chosen at program start
 - String terminator '\0'
- Gets overwritten during buffer overflow
- Check canary before jumping to return address
- Automatically inserted by compiler
 - GCC: -fstack-protector or -fstack-protector-strong

bottom of memory	DDDDDDDDDEEEEEEEEEEEEEEE	EEEE	FFFF	FFFF	FFFF	FFFF	top of memory
	89ABCDEF0123456789AB	CDEF	0123	4567	89AB	CDEF	
	buffer	sfp	cnry	ret	a	b	

<----- [] [] [] [] [] []

top of stack bottom of stack

bottom of memory	DDDDDDDDDEEEEEEEEEEEEEEE	EEEE	FFFF	FFFF	FFFF	FFFF	top of memory
	89ABCDEF0123456789AB	CDEF	0123	4567	89AB	CDEF	
	buffer	sfp	cnry	ret	a	b	

```
<----- [ NNNNNNNNNSSSSSSSSSSSS ] [ SSSS ] [ 0xDE ] [ 0xDE ] [ 0xDE ] [ 0xDE ]
```

top of
stack

bottom of
stack

Stack Canary

- **Plus:** No code changes required, only recompile
- **Minus:**
- **Minus:**
- **Minus:**


Stack Canary

- **Plus:** No code changes required, only recompile
- **Minus:** Performance penalty per return
- **Minus:** Only protects against stack smashing
- **Minus:** Fails if attacker can read memory

Buffer Overflow Defenses

- Avoid unsafe functions
- Stack canary
- Separate control stack
- Address Space Layout Randomization (ASLR)
- Memory writable or executable, not both (W^X)
- Control flow integrity (CFI)

Separate Stack




“SafeStack is an instrumentation pass that protects programs against attacks based on stack buffer overflows, without introducing any measurable performance overhead. It works by separating the program stack into two distinct regions: the safe stack and the unsafe stack. The safe stack stores **return addresses, register spills, and local variables that are always accessed in a safe way**, while the unsafe stack stores everything else. This separation ensures that buffer overflows on the unsafe stack cannot be used to overwrite anything on the safe stack.”

WebAssembly has separate stack (kind of)!


Address Space Layout Randomization

- Change location of stack, heap, code, static vars
- Works because attacker needs address of shellcode
- Layout must be unknown to attacker
 - Randomize on every launch (best)
 - Randomize at compile time
- Implemented on most modern OSes in some form

Traditional Memory Layout



PaX Memory Layout



Address Space Layout Randomization

- **Plus:** No code changes or recompile required
- **Minus:** 32-bit arch get limited protection
- **Minus:**
- **Minus:**
- **Minus:**

Address Space Layout Randomization

- **Plus:** No code changes or recompile required
- **Minus:** 32-bit arch get limited protection
- **Minus:** Fails if attacker can read memory
- **Minus:** Load-time overhead
- **Minus:** No exec img sharing between processes

W^X : write XOR execute

- Use MMU to ensure memory cannot be both writeable and executable at same time
- Code segment: executable, not writeable
- Stack, heap, static vars: writeable, not executable
- Supported by most modern processors
- Implemented by modern operating systems

W^X : write XOR execute

- **Plus:** No code changes or recompile required
- **Minus:** Requires hardware support
- **Minus:**
- **Minus:**

W^X : write XOR execute

- **Plus:** No code changes or recompile required
- **Minus:** Requires hardware support
- **Minus:** Defeated by return-oriented programming
- **Minus:** Does not protect JITed code

Buffer Overflow Defenses

- Avoid unsafe functions
- Stack canary
- Separate control stack
- Address Space Layout Randomization (ASLR)
- Memory writable or executable, not both (W^X)
- Control flow integrity (CFI)

Control Flow Integrity

- Check destination of every indirect jump
 - Function returns
 - Function pointers
 - Virtual methods
- What are the valid destinations?
 -
 -

Control Flow Integrity

- Check destination of every indirect jump
 - Function returns
 - Function pointers
 - Virtual methods
- What are the valid destinations?
 - Caller of every function known at compile time
 - Class hierarchy limits possible virtual function instances

CFI

- **Plus:** No code changes or hardware support
- **Plus:** Protects against many vulnerabilities
- **Minus:**
- **Minus:**
- **Minus:**

CFI

- **Plus:** No code changes or hardware support
- **Plus:** Protects against many vulnerabilities
- **Minus:** Performance overhead
- **Minus:** Requires smarter compiler
- **Minus:** Requires having all code available